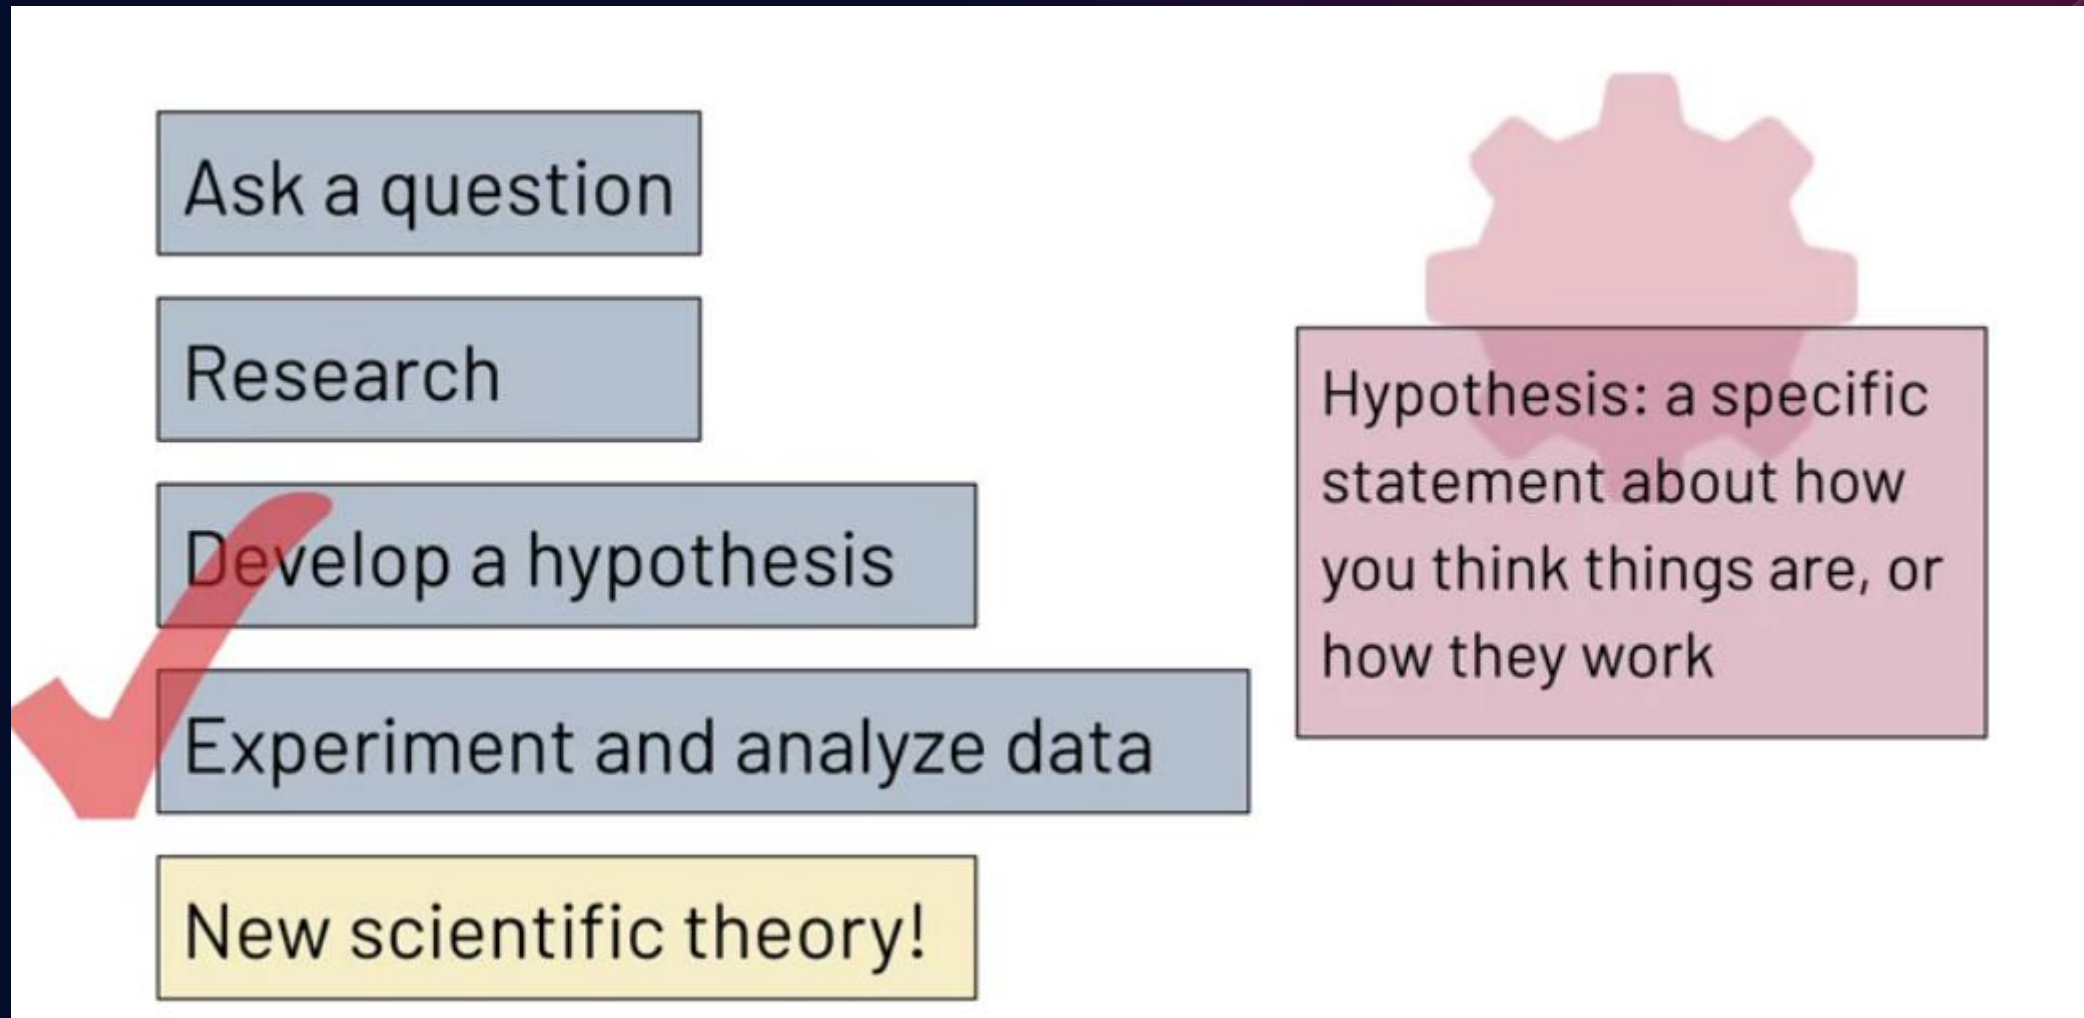


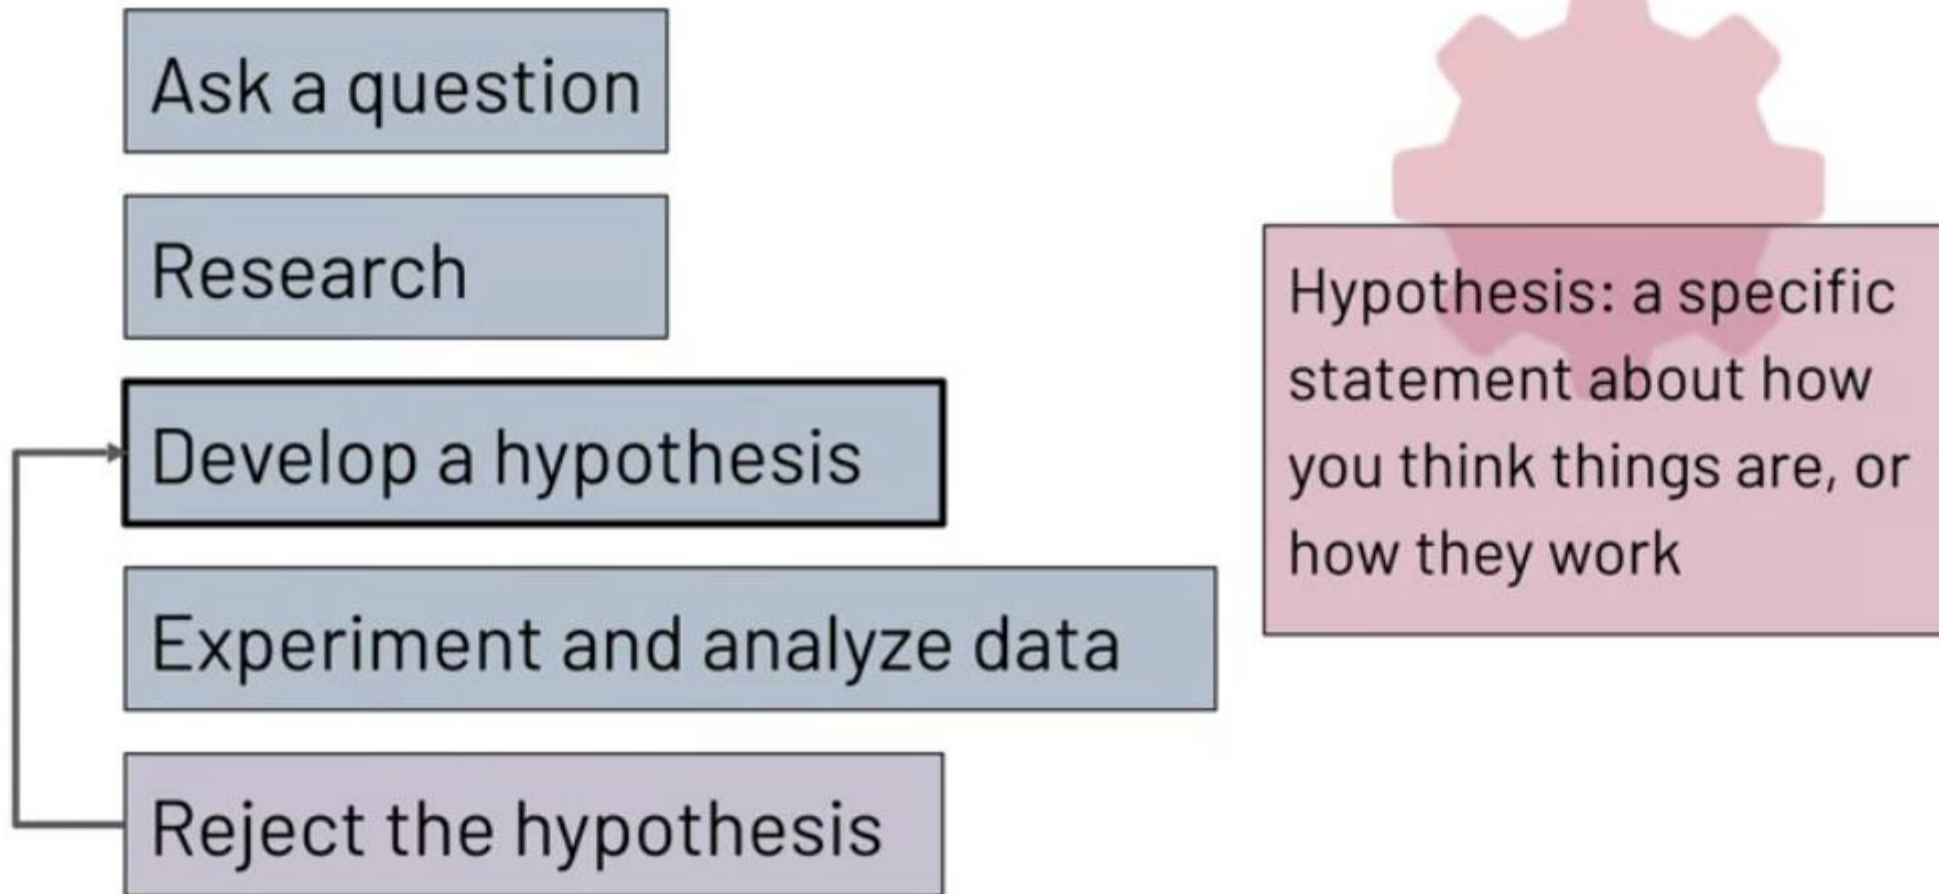
# CS 156: Introduction to Artificial Intelligence

**Instructor: Dr. Sayma Akther**  
San José State University

# Classification Models in Machine Learning

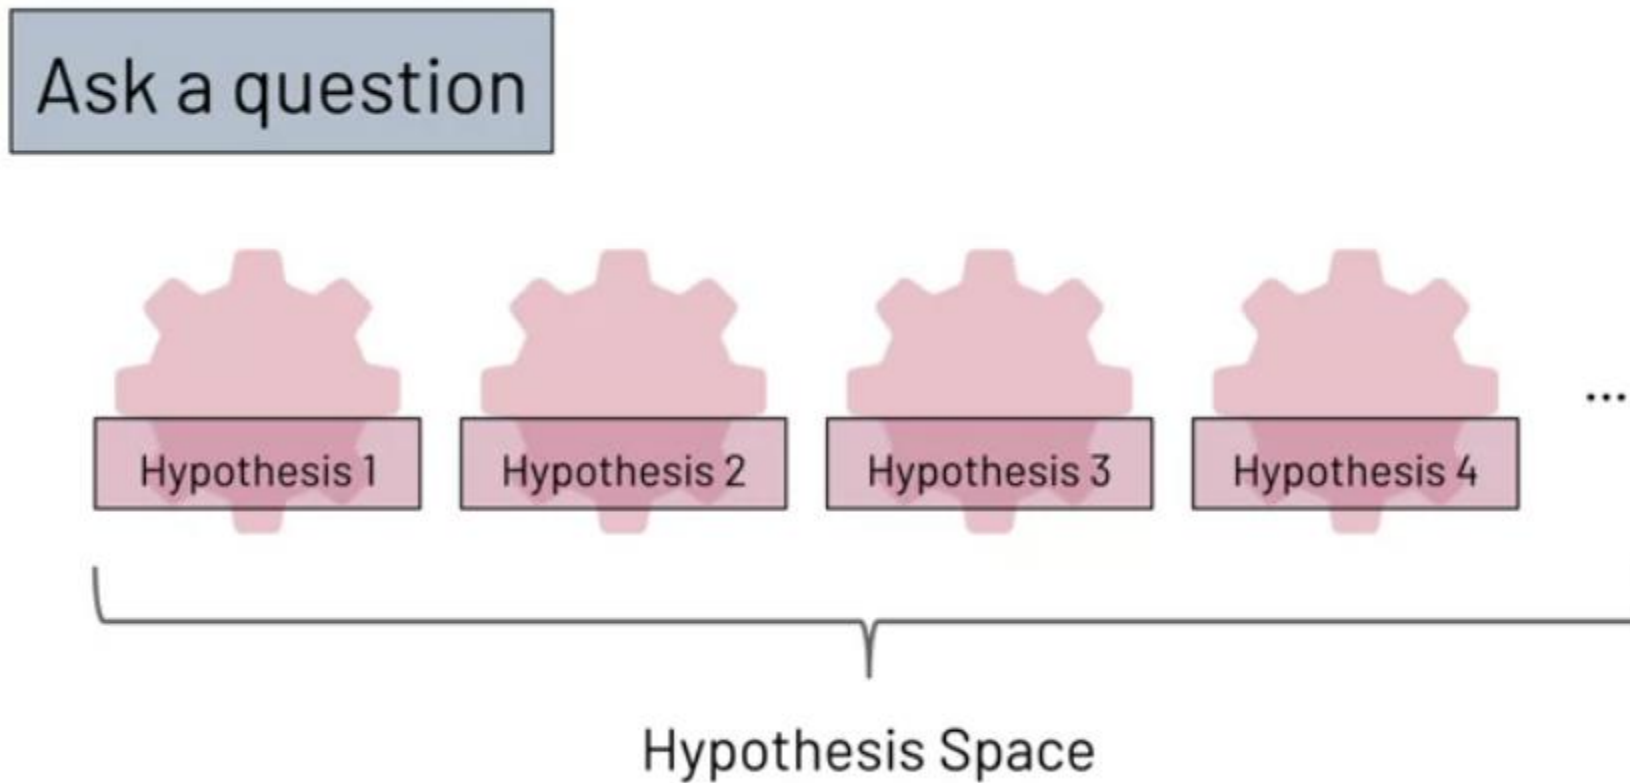


# Classification Models in Machine Learning



slide~superdatascience

# Classification Models in Machine Learning

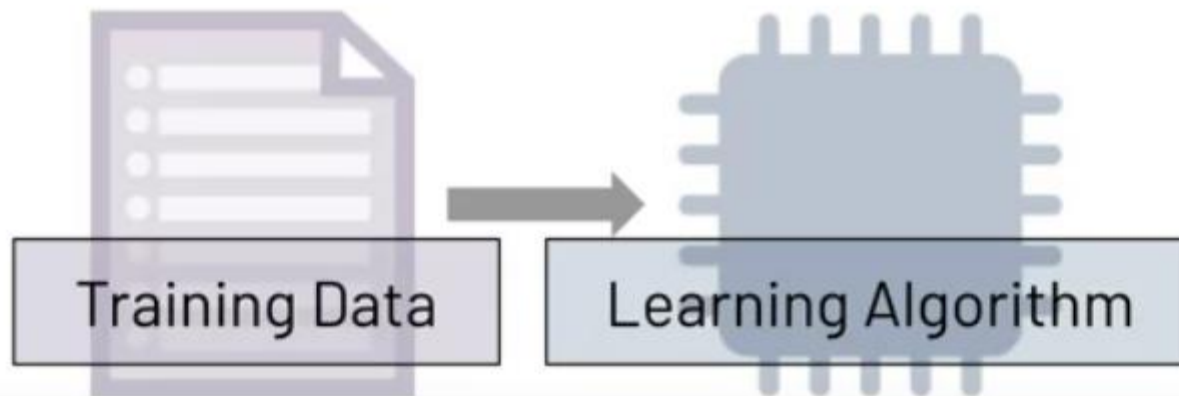




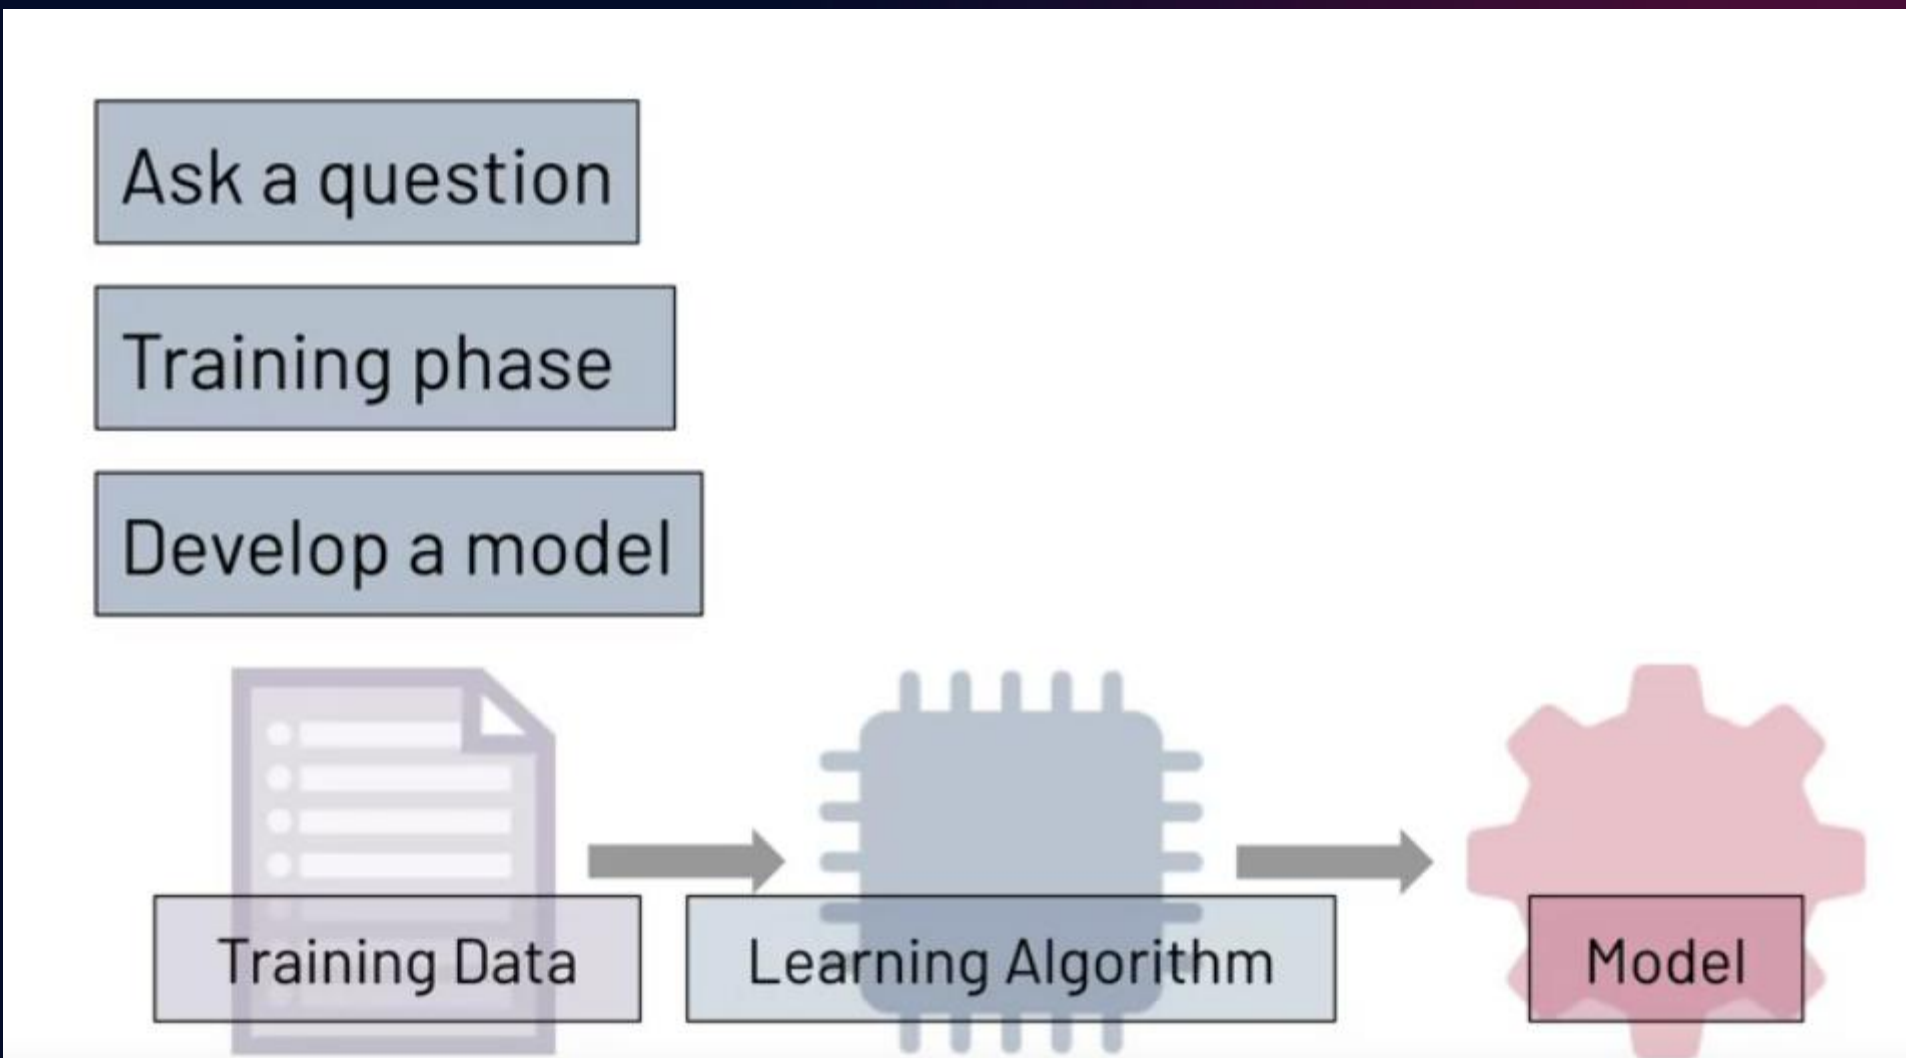
# Classification Models in Machine Learning

Ask a question

Training phase

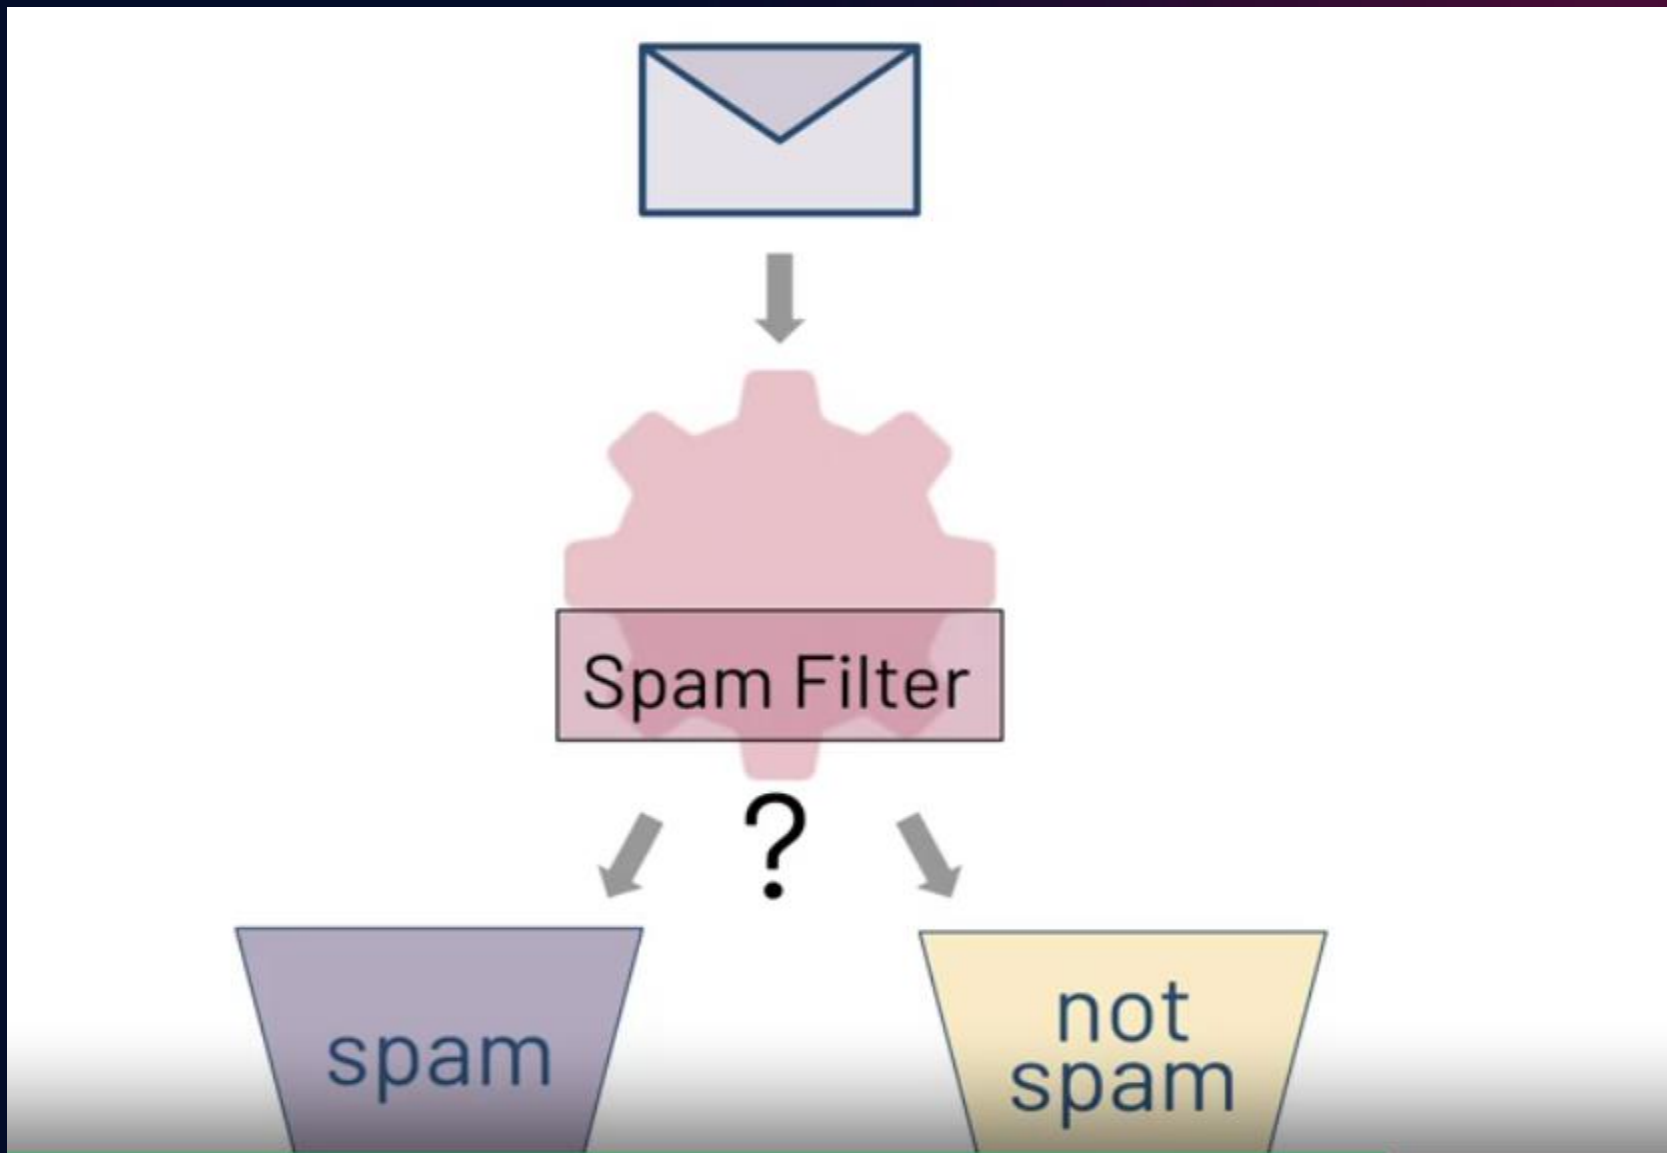


# Classification Models in Machine Learning



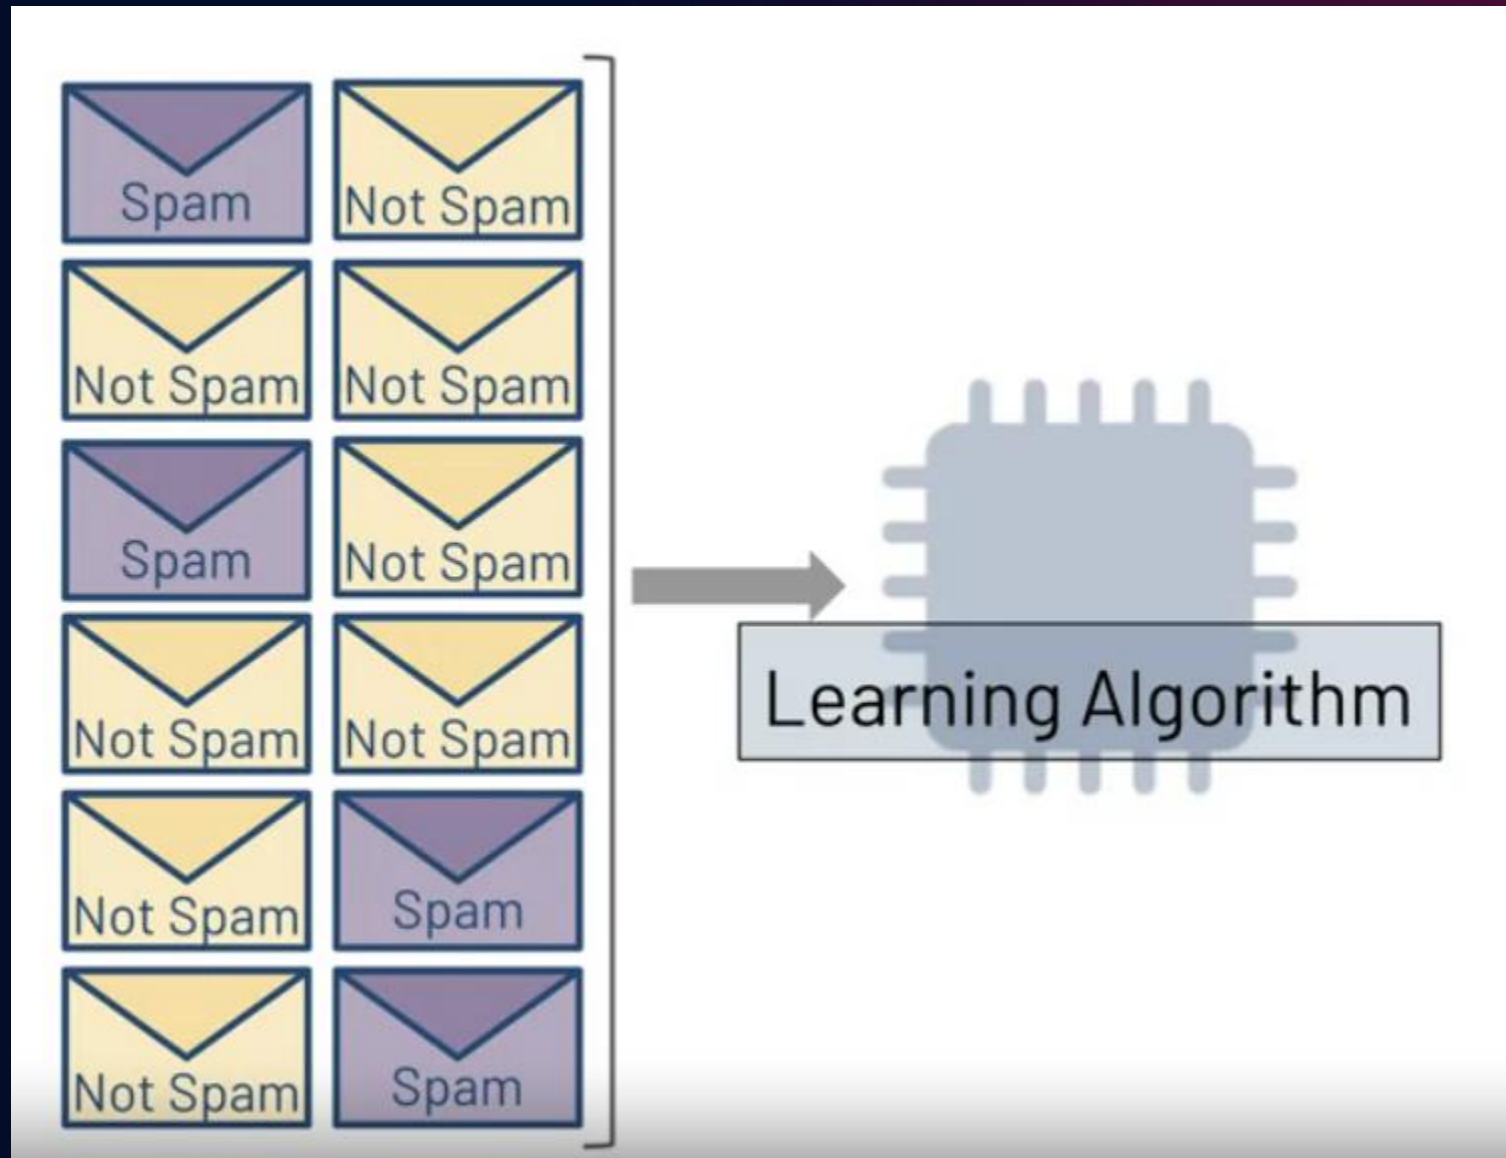


# Classification Models in Machine Learning



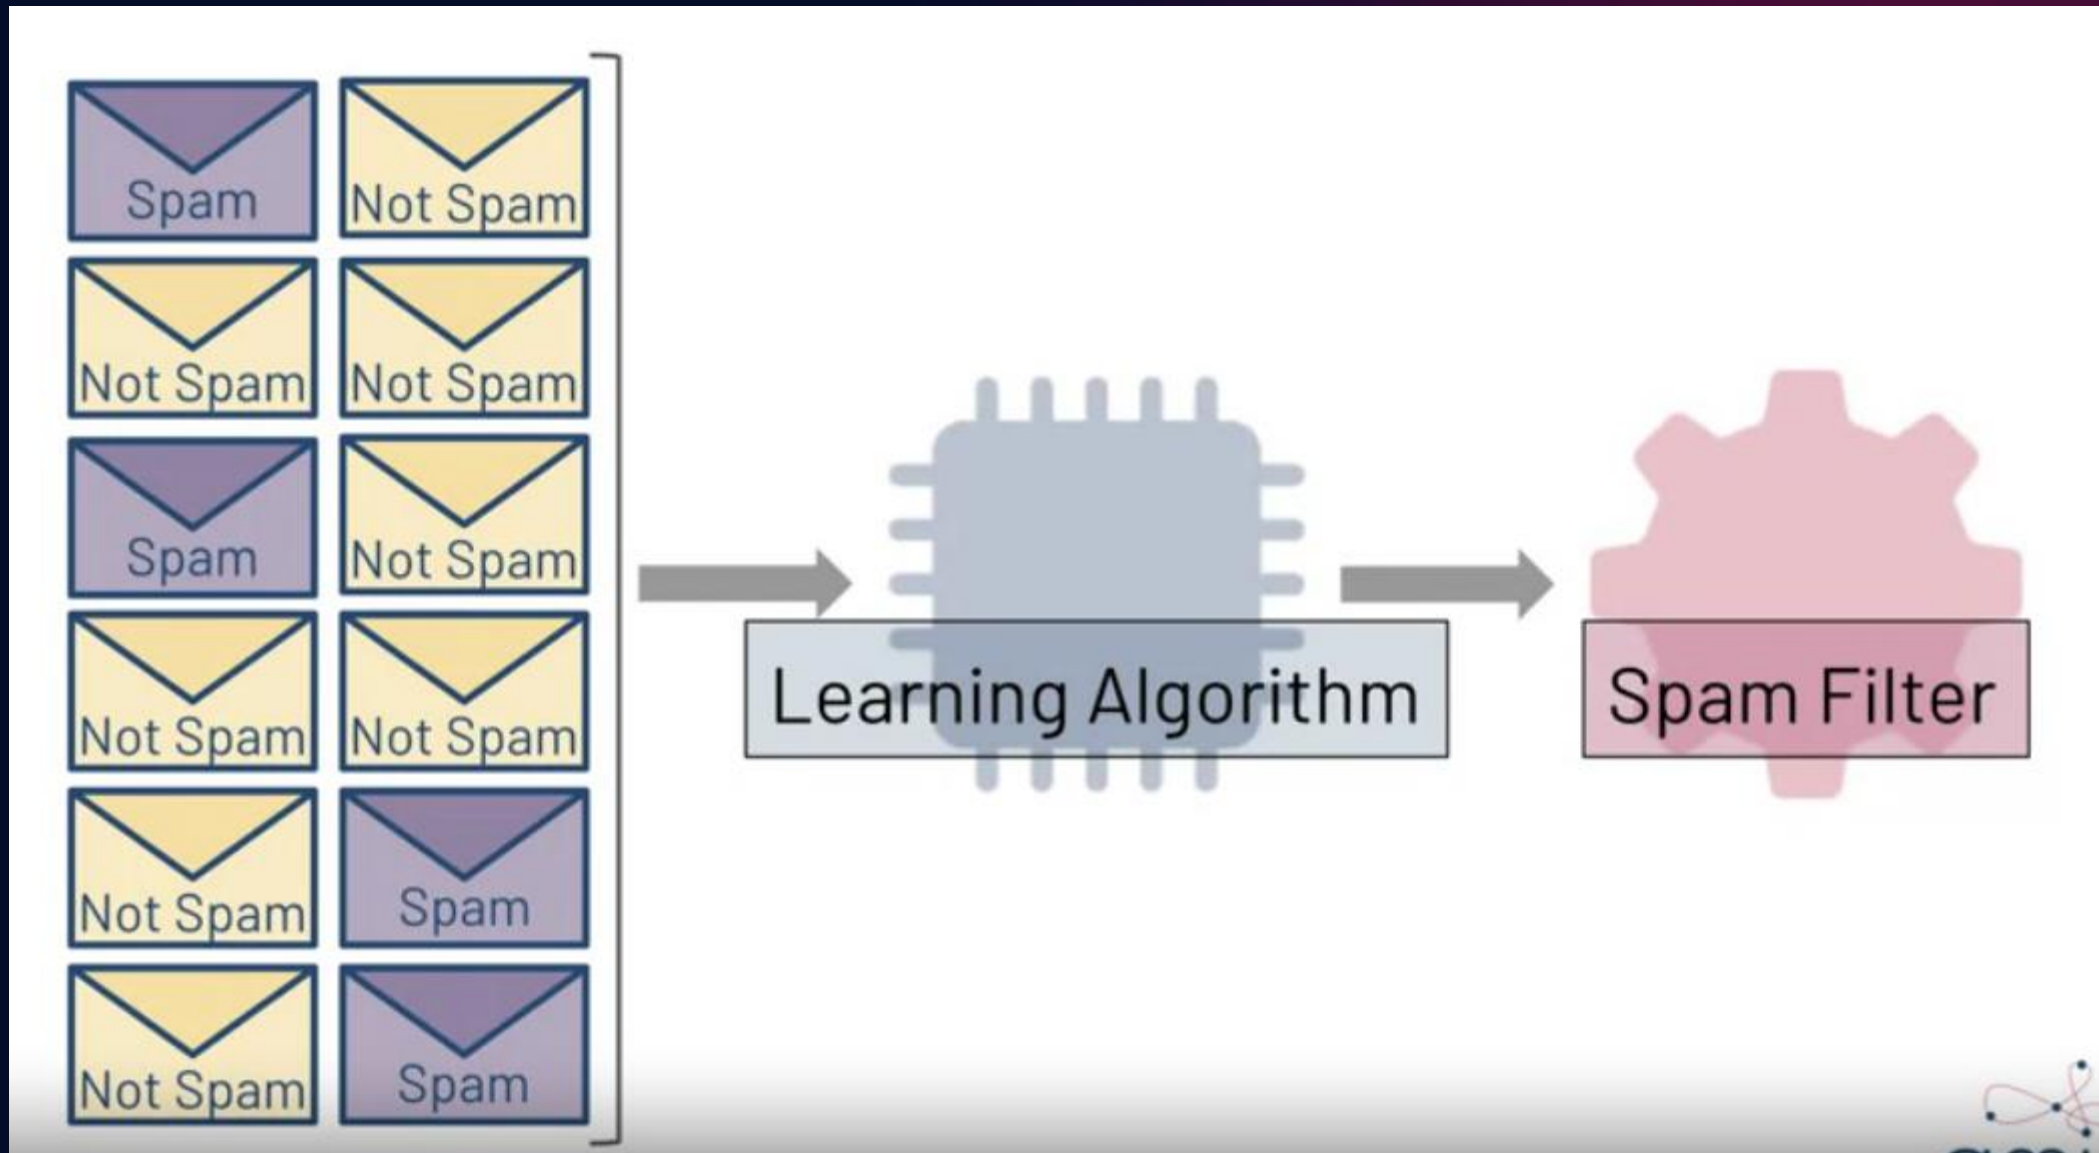


# Classification Models in Machine Learning





# Classification Models in Machine Learning

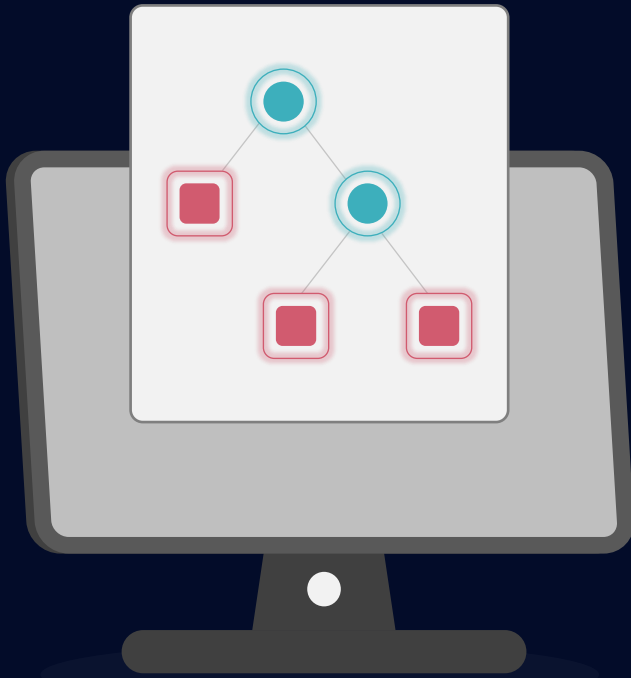


# Decision Tree

It is a handy tool with many applications. Decision trees can be used to solve classification and regression issues

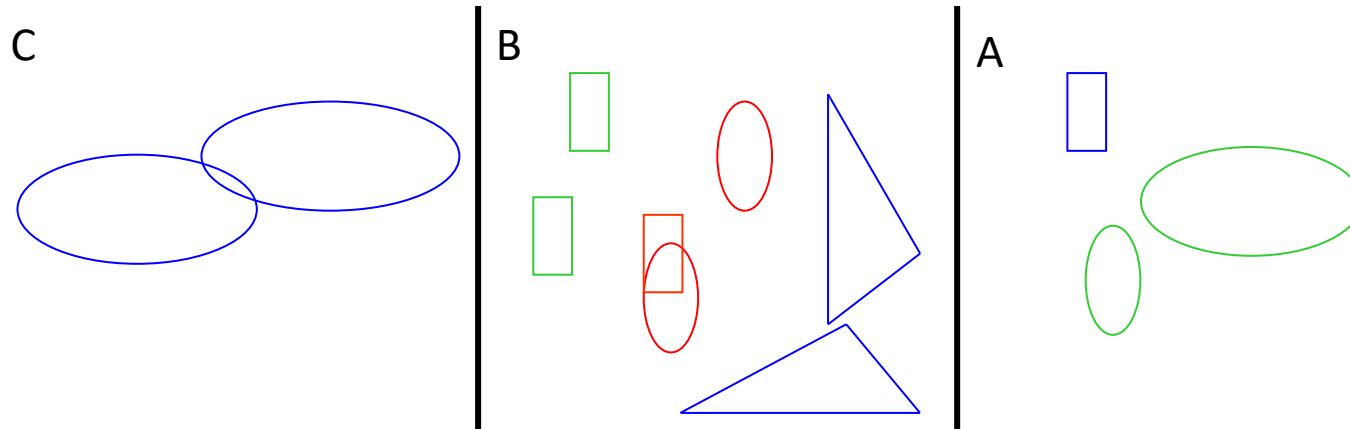
The name indicates that it displays the predictions coming from a series of feature-based splits using a flowchart-like tree structure

It all starts with a root node and ends with a leaf choice



# Decision Trees

- A hierarchical data structure that represents data by implementing a divide and conquer strategy
- Can be used as a non-parametric classification and regression method
- Given a collection of examples, learn a decision tree that represents it.
- Use this representation to classify new examples

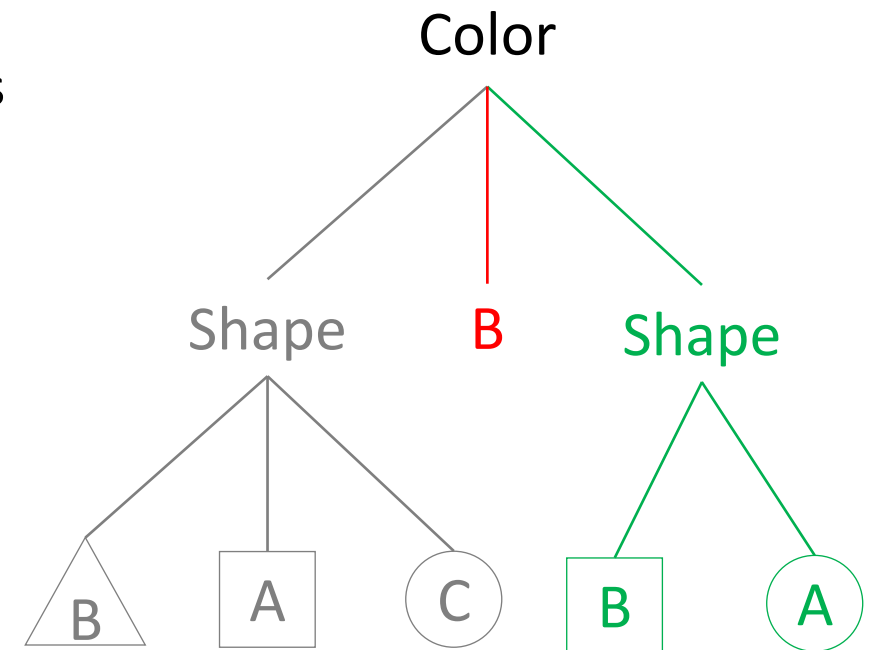
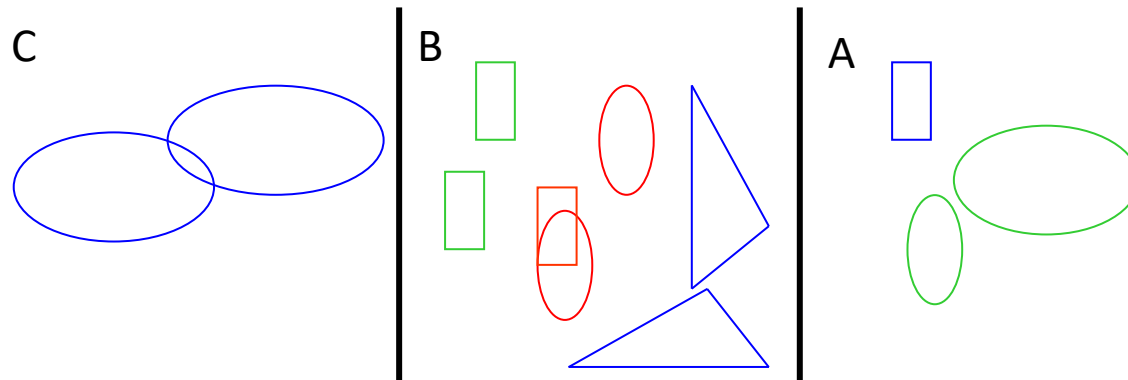


# The Representation

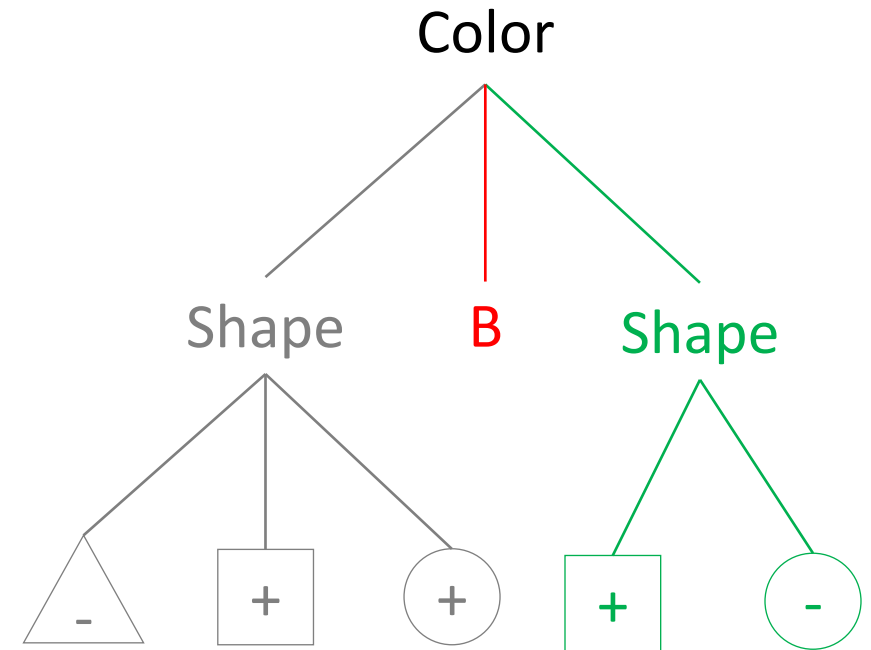
- Decision Trees are classifiers for instances represented as feature vectors
  - $\text{color}=\{\text{red, blue, green}\}$  ;  $\text{shape}=\{\text{circle, triangle, rectangle}\}$  ;  $\text{label}=\{A, B, C\}$
- **Nodes** are **tests** for feature values
- There is one branch for each value of the feature
- **Leaves** specify the category (labels)
- Can categorize instances into multiple disjoint categories

Evaluation of a  
Decision Tree

Learning a  
Decision Tree

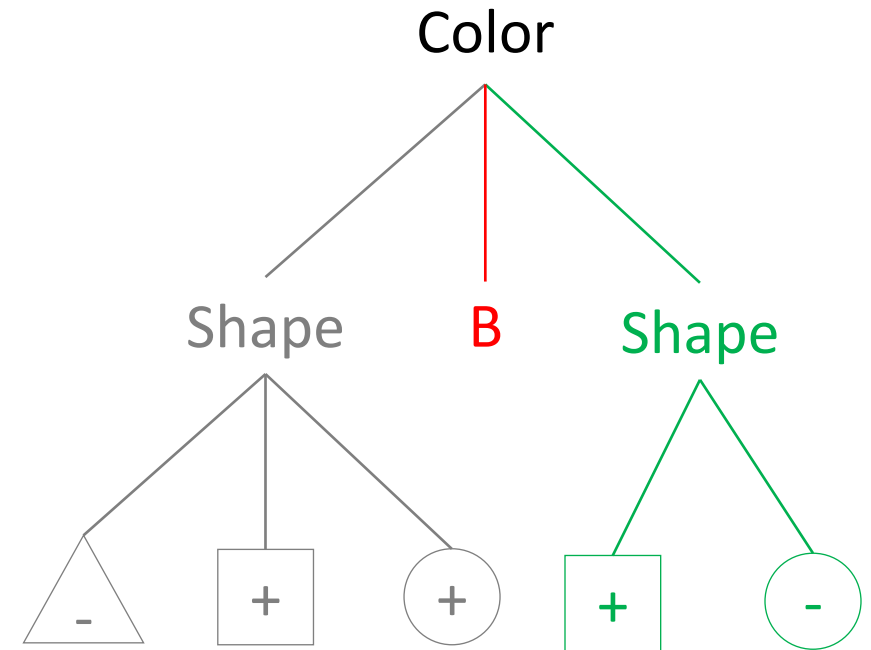


# Expressivity of Decision Trees



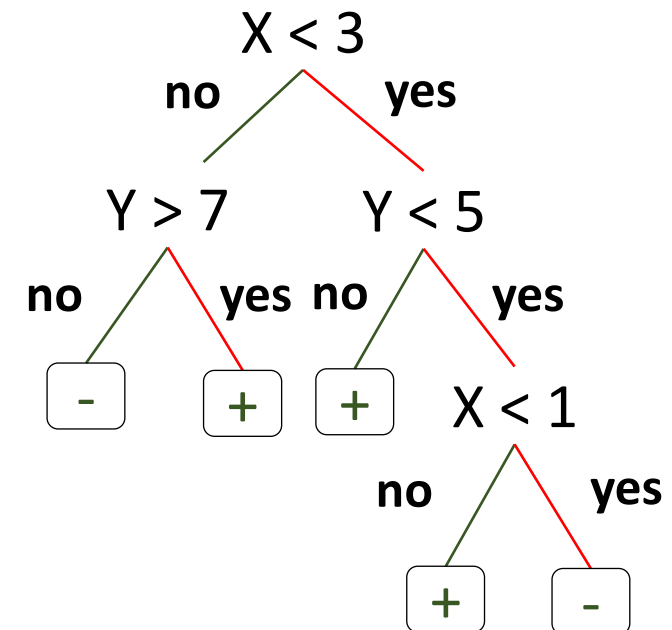
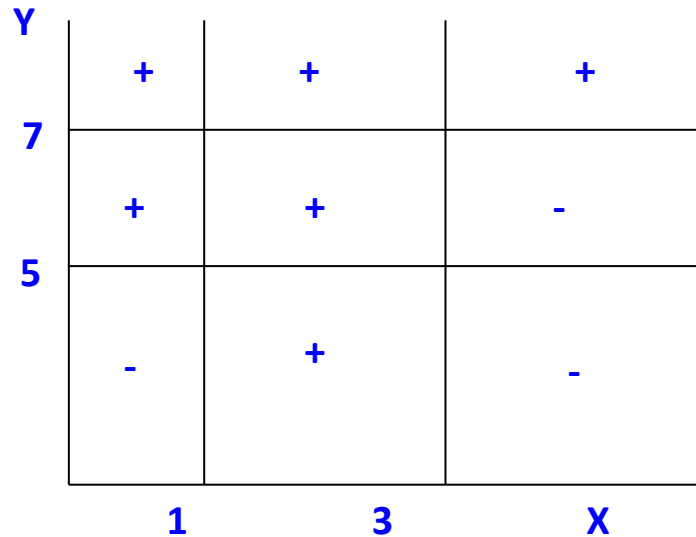
# Decision Trees

- Output is a discrete category. Real valued outputs are possible (regression trees)
- There are efficient algorithms for processing large amounts of data (but not too many features)
- There are methods for handling **noisy data** (classification noise and attribute noise) and for handling missing attribute values



# Decision Boundaries

- Usually, instances are represented as attribute-value pairs (color=blue, shape = square, +)
- Numerical values can be used either by discretizing or by using thresholds for splitting nodes
- In this case, the tree divides the features space into axis-parallel rectangles, each labeled with one of the labels





# Today's key concepts

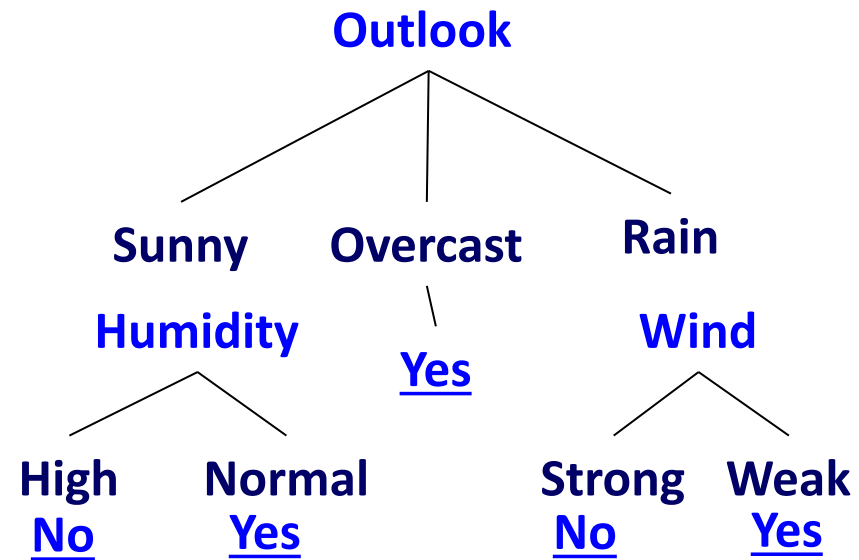
---

- Learning decision trees (ID3 algorithm)
  - Greedy heuristic (based on information gain)  
Originally developed for discrete features

# Learning decision trees (ID3 algorithm

# Decision Trees

- Can represent any Boolean Function
- Can be viewed as a way to compactly represent a lot of data.
- Natural representation: (20 questions)
- The **evaluation** of the Decision Tree Classifier is easy
- Clearly, given data, there are many ways to represent it as a decision tree.
- Learning a **good** representation from data is the challenge.



# Will I play tennis today?

---

- **Features**

- Outlook: {Sun, Overcast, Rain}
- Temperature: {Hot, Mild, Cool}
- Humidity: {High, Normal, Low}
- Wind: {Strong, Weak}

- **Labels**

- Binary classification task:  $Y = \{+, -\}$

# Will I play tennis today?

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

**Outlook:** S(unny),  
O(vercast),  
R(ainy)

**Temperature:** H(ot),  
M(edium),  
C(ool)

**Humidity:** H(igh),  
N(ormal),  
L(ow)

**Wind:** S(trong),  
W(eak)

# Basic Decision Trees Learning Algorithm

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

Algorithm?

# LEARN-DECISION-TREE

```
function LEARN-DECISION-TREE(examples, attributes, parent_examples) returns a tree

if examples is empty then return PLURALITY-VALUE(parent_examples)
else if all examples have the same classification then return the classification
else if attributes is empty then return PLURALITY-VALUE(examples)
else
     $A \leftarrow \operatorname{argmax}_{a \in \text{attributes}} \text{IMPORTANCE}(a, \text{examples})$ 
    tree  $\leftarrow$  a new decision tree with root test A
    for each value v of A do
        exs  $\leftarrow \{e : e \in \text{examples} \text{ and } e.A = v\}$ 
        subtree  $\leftarrow$  LEARN-DECISION-TREE(exs, attributes − A, examples)
        add a branch to tree with label (A = v) and subtree subtree
    return tree
```

The function PLURALITY-VALUE selects the most common output value among a set of examples, breaking ties randomly.



# Picking the Root Attribute

---

- The goal is to have the resulting decision tree as small as possible (Occam's Razor)
  - The main decision in the algorithm is the selection of the next attribute to condition on.
- We want attributes that split the examples to sets that are **relatively pure in one label**; this way we are closer to a leaf node.
  - The most popular heuristics is based on **information gain**, originated with the ID3 system of Quinlan.

# Entropy

- Entropy (impurity, disorder) of a set of examples,  $S$ , relative to a binary classification is:

$$\text{Entropy}(S) = -p_+ \log(p_+) - p_- \log(p_-)$$

- $p_+$  is the proportion of positive examples in  $S$  and
- $p_-$  is the proportion of negative examples in  $S$ 
  - If all the examples belong to the same category: Entropy = 0
  - If all the examples are equally mixed (0.5, 0.5): Entropy = 1
  - Entropy = Level of uncertainty.

- In general, when  $p_i$  is the fraction of examples labeled  $i$ :

$$\text{Entropy}(S[p_1, p_2, \dots, p_k]) = -\sum_{i=1}^k p_i \log(p_i)$$

- Entropy can be viewed as the number of bits required, on average, to encode the class of labels. If the probability for  $+$  is 0.5, a single bit is required for each example; if it is 0.8 – can use less than 1 bit.

# Information Gain

High Entropy – High level of Uncertainty

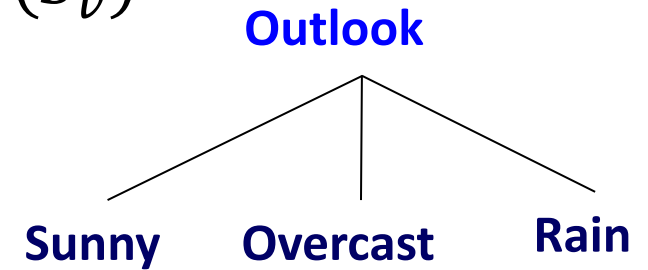
Low Entropy – No Uncertainty.

- The information gain of an attribute **a** is the expected reduction in entropy caused by partitioning on this attribute

$$Gain(S, a) = Entropy(S) - \sum_{v \in values(S)} \frac{|S_v|}{|S|} Entropy(S_v)$$

- Where:

- $S_v$  is the subset of **S** for which attribute **a** has value **v**, and
- the entropy of partitioning the data is calculated by **weighing the entropy of each partition** by its size relative to the original set



- Partitions of low entropy (imbalanced splits) lead to high gain
- Go back to check which of the A, B splits is better

# Will I play tennis today?

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

**Outlook:** S(unny),  
O(vercast),  
R(ainy)

**Temperature:** H(ot),  
M(edium),  
C(ool)

**Humidity:** H(igh),  
N(ormal),  
L(ow)

**Wind:** S(trong),  
W(eak)

# Will I play tennis today?

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

calculate current entropy

- $p_+ = \frac{9}{14}$     $p_- = \frac{5}{14}$

- $Entropy(Play) = -p_+ \log_2(p_+) - p_- \log_2(p_-)$   
 $= -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14}$   
 $\approx 0.94$

# Information Gain: Outlook

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

$$Gain(S, a) = Entropy(S) - \sum_{v \in values(S)} \frac{|S_v|}{|S|} Entropy(S_v)$$

**Outlook = sunny:**

$$p_+ = 2/5 \quad p_- = 3/5 \quad Entropy(O = S) = 0.971$$

**Outlook = overcast:**

$$p_+ = 4/4 \quad p_- = 0 \quad Entropy(O = O) = 0$$

**Outlook = rainy:**

$$p_+ = 3/5 \quad p_- = 2/5 \quad Entropy(O = R) = 0.971$$

**Expected entropy**

$$= \sum_{v \in values(S)} \frac{|S_v|}{|S|} Entropy(S_v)$$

$$= (5/14) \times 0.971 + (4/14) \times 0 + (5/14) \times 0.971 = \mathbf{0.694}$$

$$\text{Information gain} = 0.940 - 0.694 = \mathbf{0.246}$$

# Information Gain: Humidity

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

$$Gain(S, a) = Entropy(S) - \sum_{v \in values(S)} \frac{|S_v|}{|S|} Entropy(S_v)$$

**Humidity = high:**

$$p_+ = 3/7 \quad p_- = 4/7 \quad Entropy(H = H) = 0.985$$

**Humidity = Normal:**

$$p_+ = 6/7 \quad p_- = 1/7 \quad Entropy(H = N) = 0.592$$

**Expected entropy**

$$= \sum_{v \in values(S)} \frac{|S_v|}{|S|} Entropy(S_v)$$

$$= (7/14) \times 0.985 + (7/14) \times 0.592 = \mathbf{0.7785}$$

$$\mathbf{Information\ gain = 0.940 - 0.7785 = 0.1615}$$



# Which feature to split on?

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

## Information gain:

Outlook: 0.246

Humidity: 0.1615

Wind:?

Temperature:?

→ Split on Outlook

# Which feature to split on?

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

## Information gain:

Outlook: 0.246

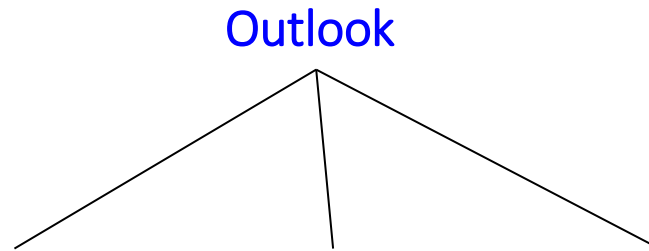
Humidity: 0.1615

Wind: 0.048

Temperature: 0.029

→ Split on Outlook

# An Illustrative Example (III)



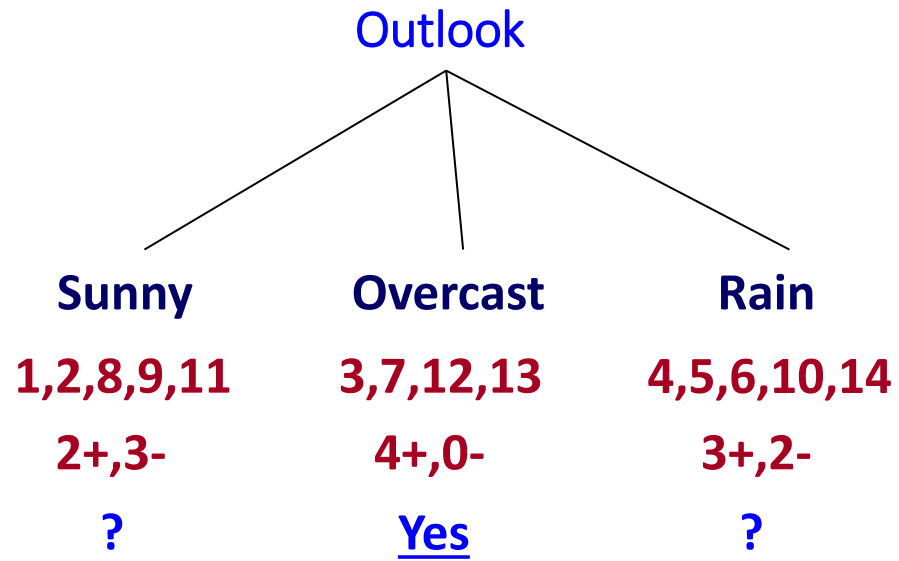
Gain(S, Humidity) = 0.1615

Gain(S, Wind) = 0.048

Gain(S, Temperature) = 0.029

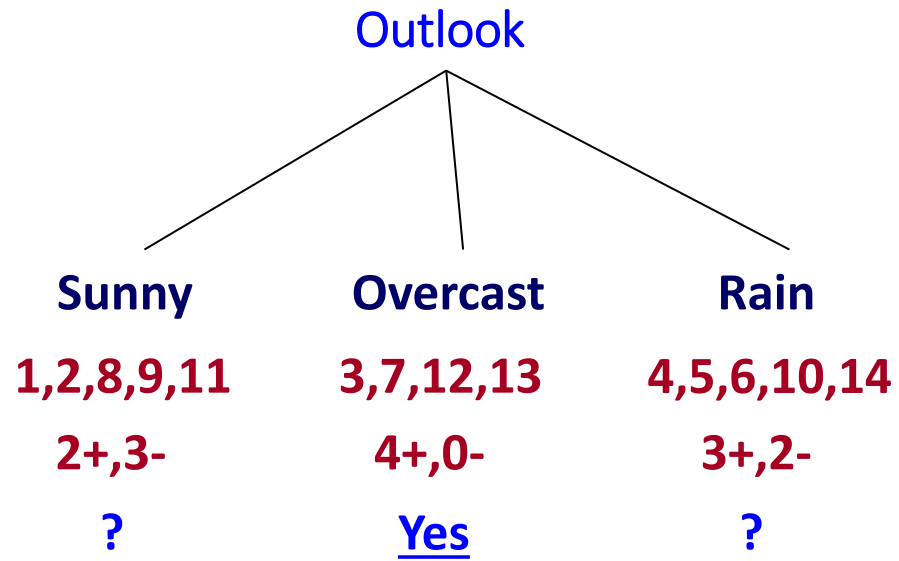
Gain(S, Outlook) = 0.246

# An Illustrative Example (III)



	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

# An Illustrative Example (III)

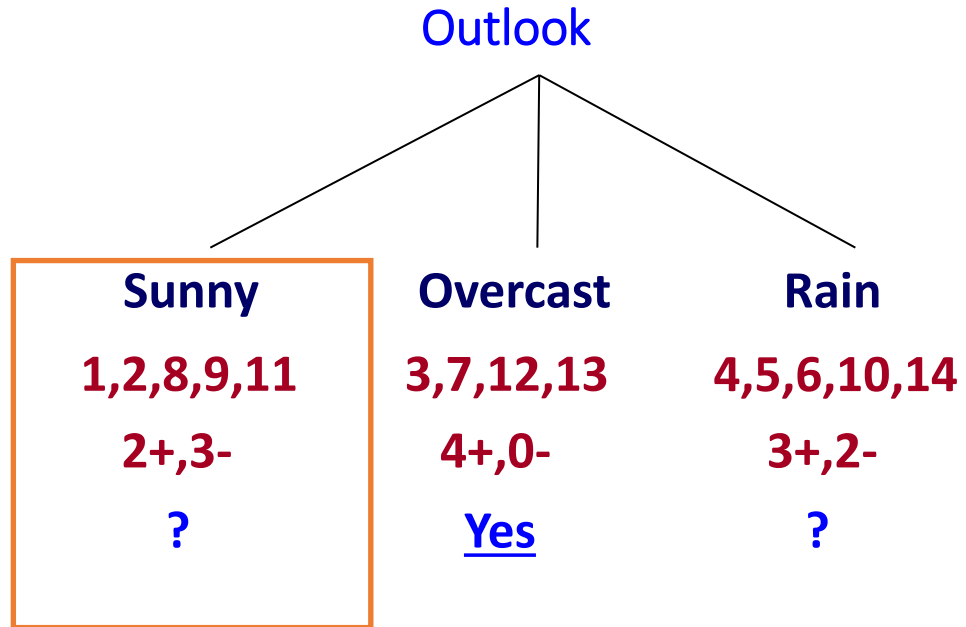


Continue until:

- Every attribute is included in **path**, or,
- All examples in the leaf have same label

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

# An Illustrative Example (IV)



$$\text{Gain}(S_{\text{sunny}}, \text{Humidity}) = .97 - (3/5) 0 - (2/5) 0 = .97$$

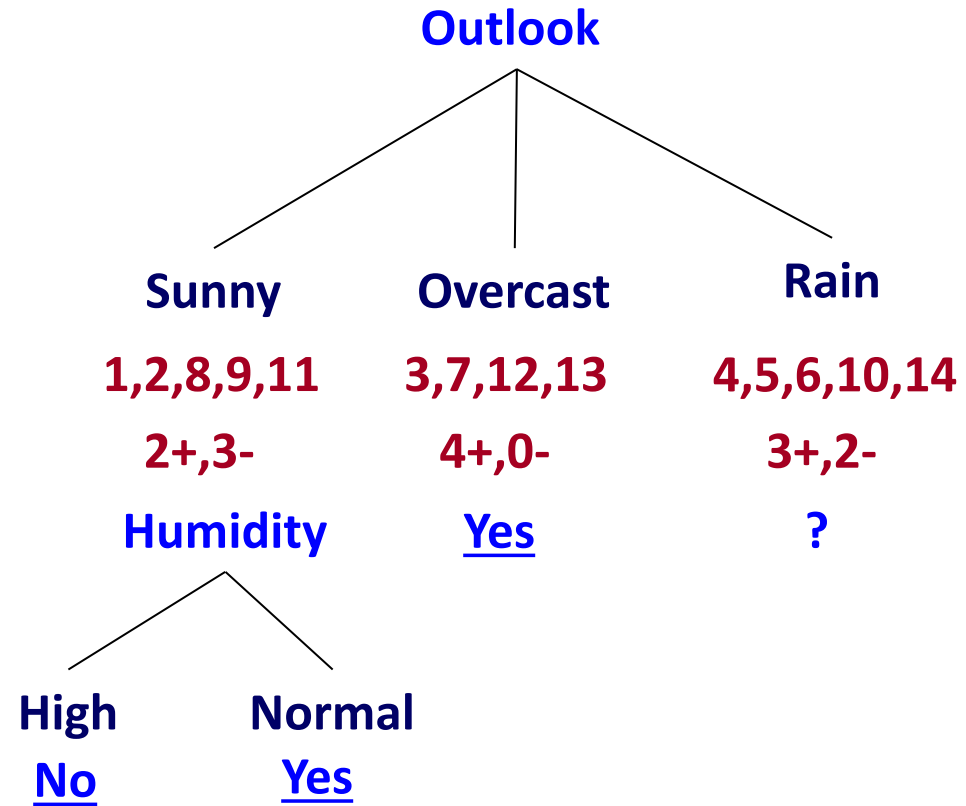
$$\text{Gain}(S_{\text{sunny}}, \text{Temp}) = .97 - 0 - (2/5) 1 = .57$$

$$\text{Gain}(S_{\text{sunny}}, \text{Wind}) = .97 - (2/5) 1 - (3/5) .92 = .02$$

Split on Humidity

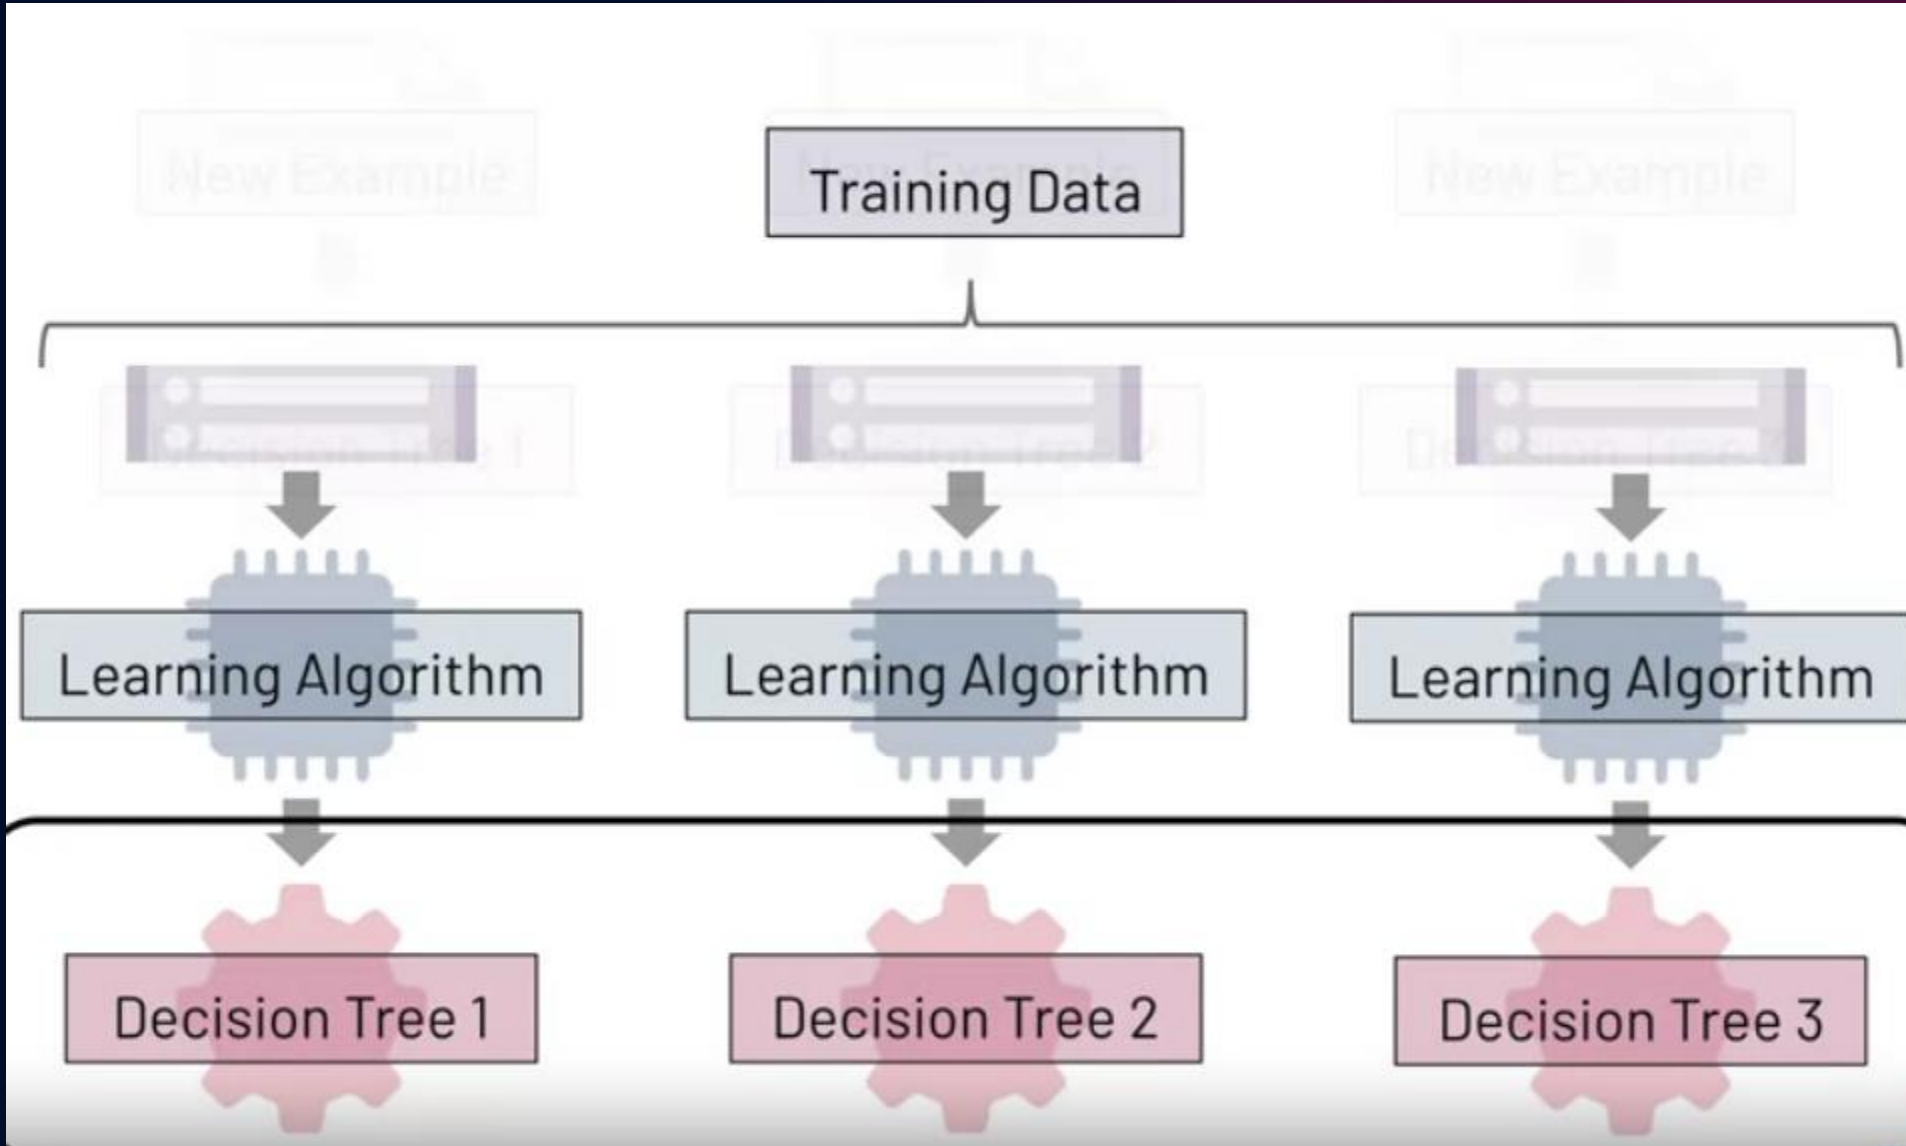
	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

# An Illustrative Example (V)





# Classification Models: Decision Tree





# Random Forest Algorithm

## Ensemble Learning



# Random Forest Algorithm

STEP 1: Pick at random  $K$  data points from the Training set.



STEP 2: Build the Decision Tree associated to these  $K$  data points.



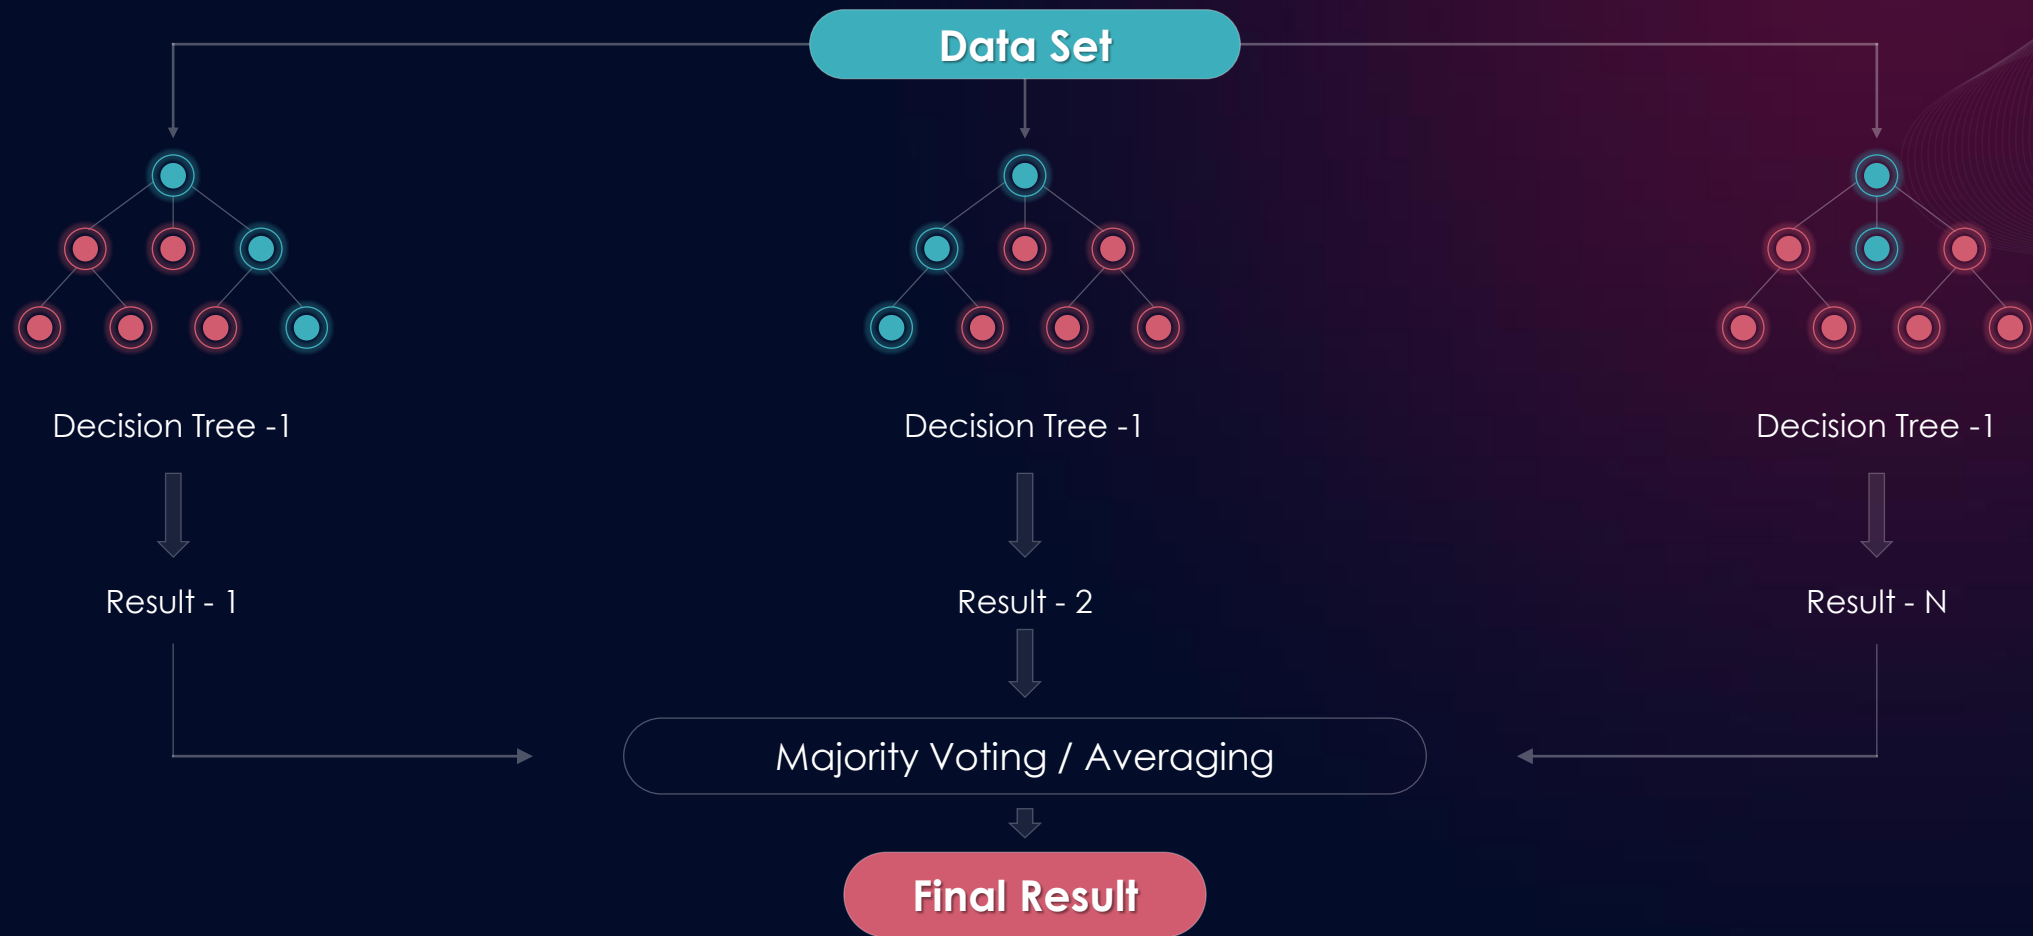
STEP 3: Choose the number  $N_{tree}$  of trees you want to build and repeat STEPS 1 & 2



STEP 4: For a new data point, make each one of your  $N_{tree}$  trees predict the value of  $Y$  for the data point in question, and assign the new data point the average across all of the predicted  $Y$  values.



# Random Forest Algorithm



A Random Forest is a cluster of decision trees. Each tree is classed, and the tree "votes" for that class to classify a new item based on its properties. The forest chooses the categorization with the highest votes (over all the trees in the forest).

# KNN (K- Nearest Neighbors) Algorithm

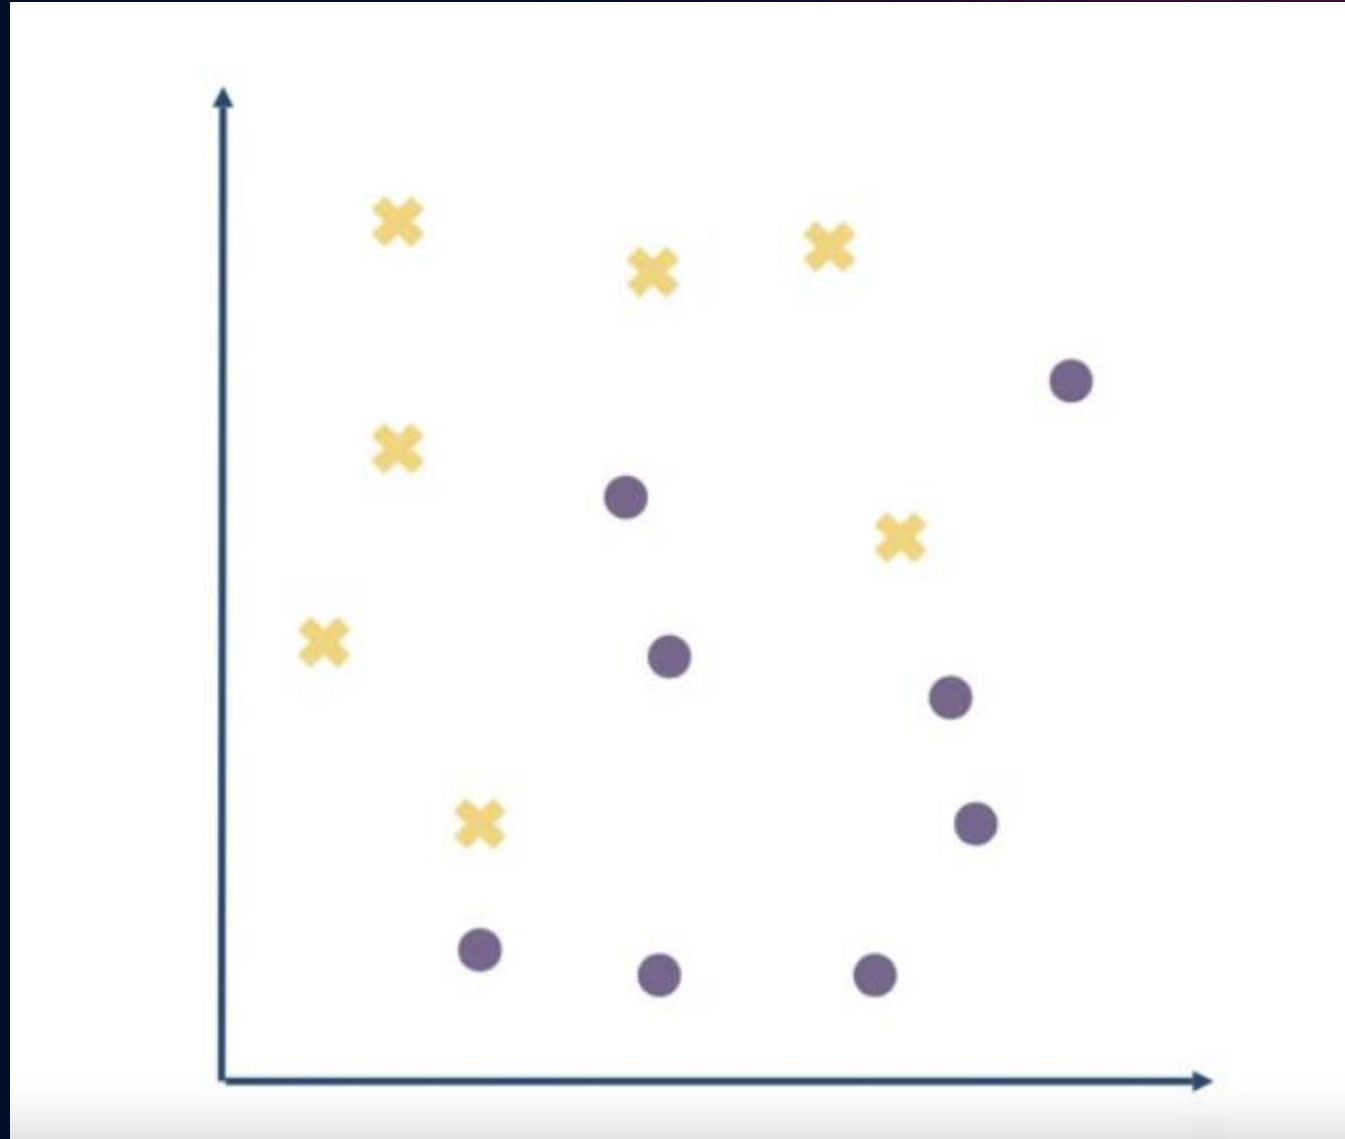


It's a simple algorithm that keeps all existing instances, and classifies new cases based on a majority vote of its  $k$  neighbors

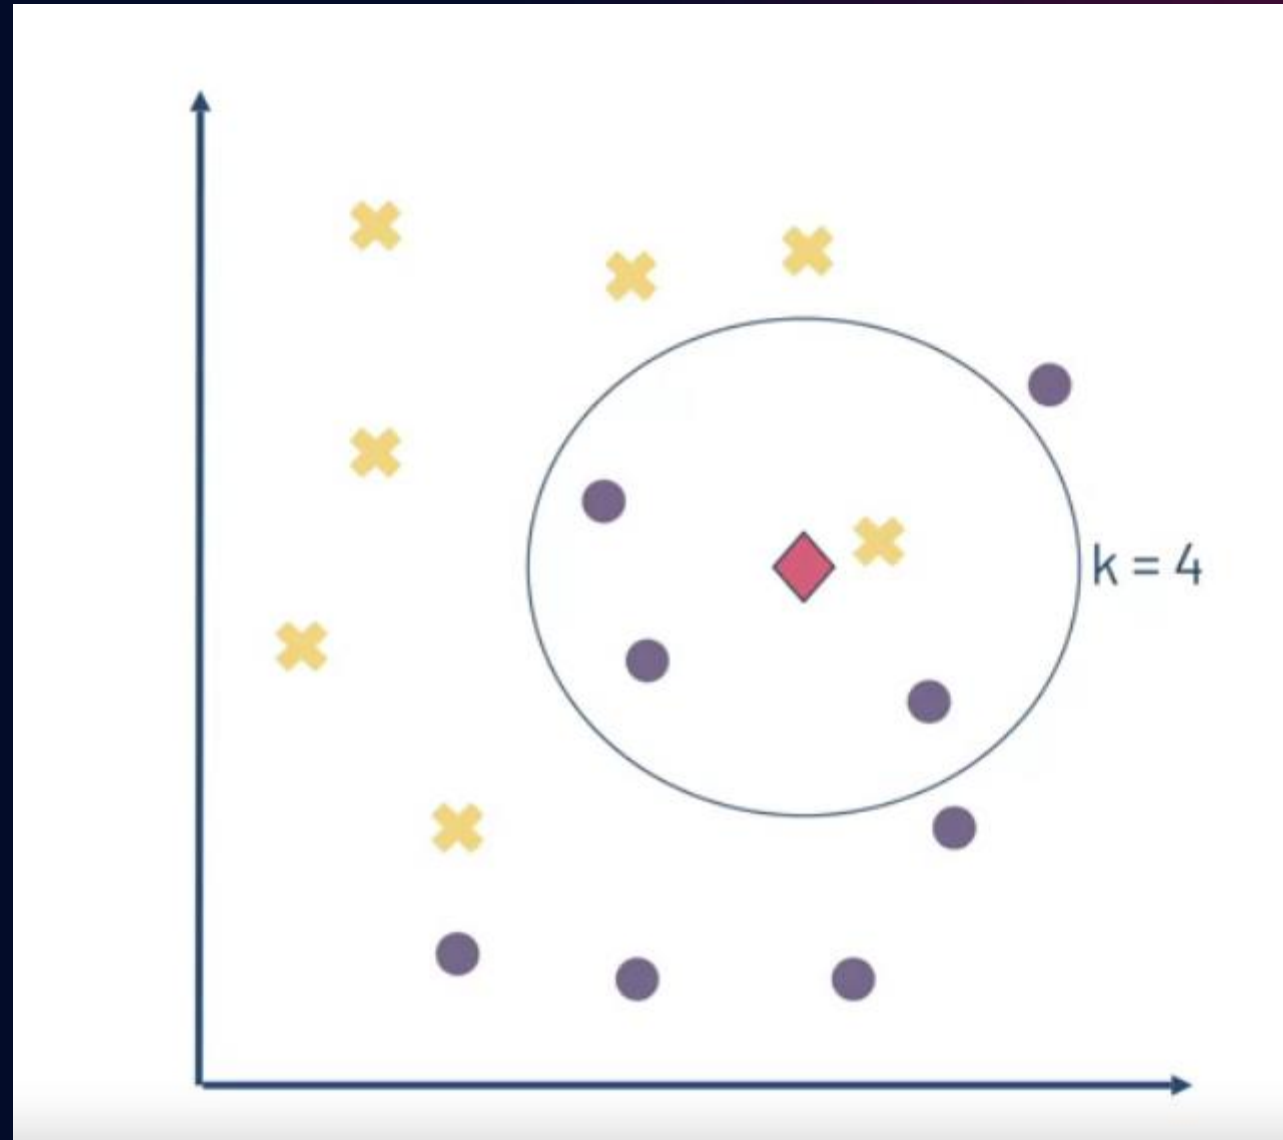


The case is then given to the class with which it has greatest in common. A distance function is used to perform this task

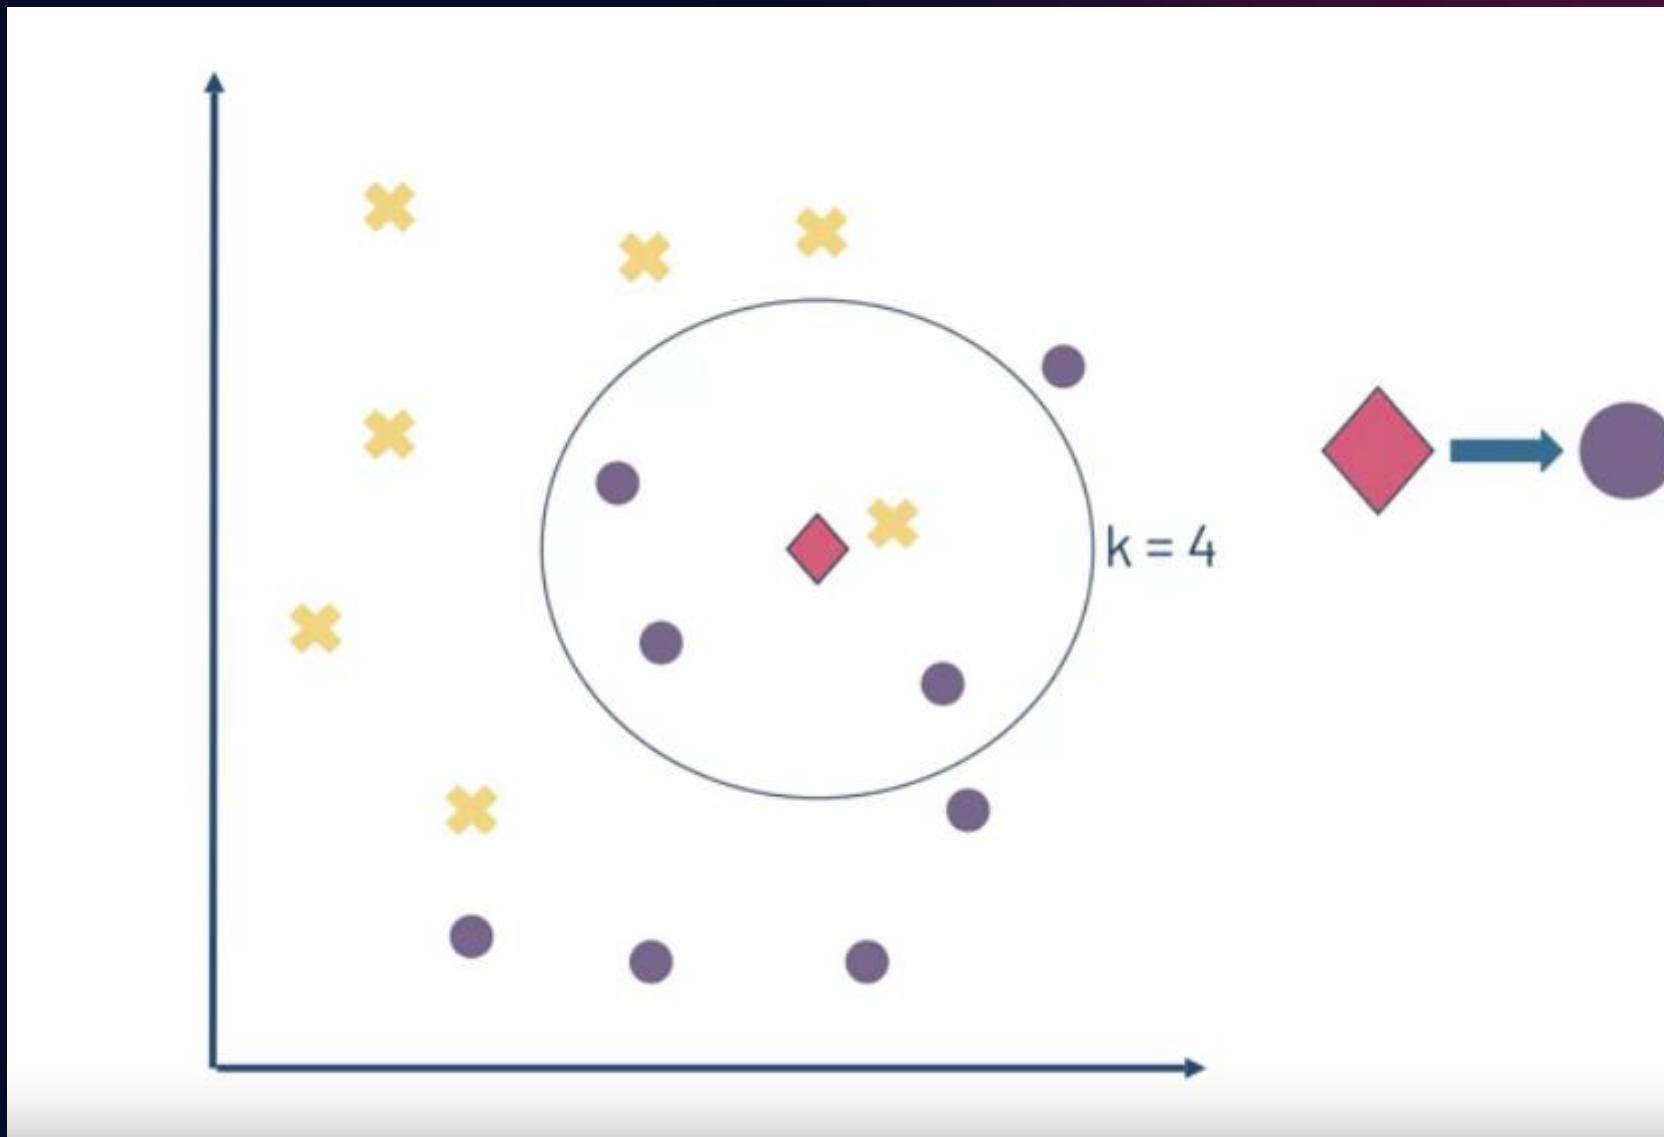
# KNN (K- Nearest Neighbors) Algorithm



# KNN (K- Nearest Neighbors) Algorithm



# KNN (K- Nearest Neighbors) Algorithm





# KNN (K- Nearest Neighbors) Algorithm






# KNN (K- Nearest Neighbors) Algorithm

[number of rooms, square footage]

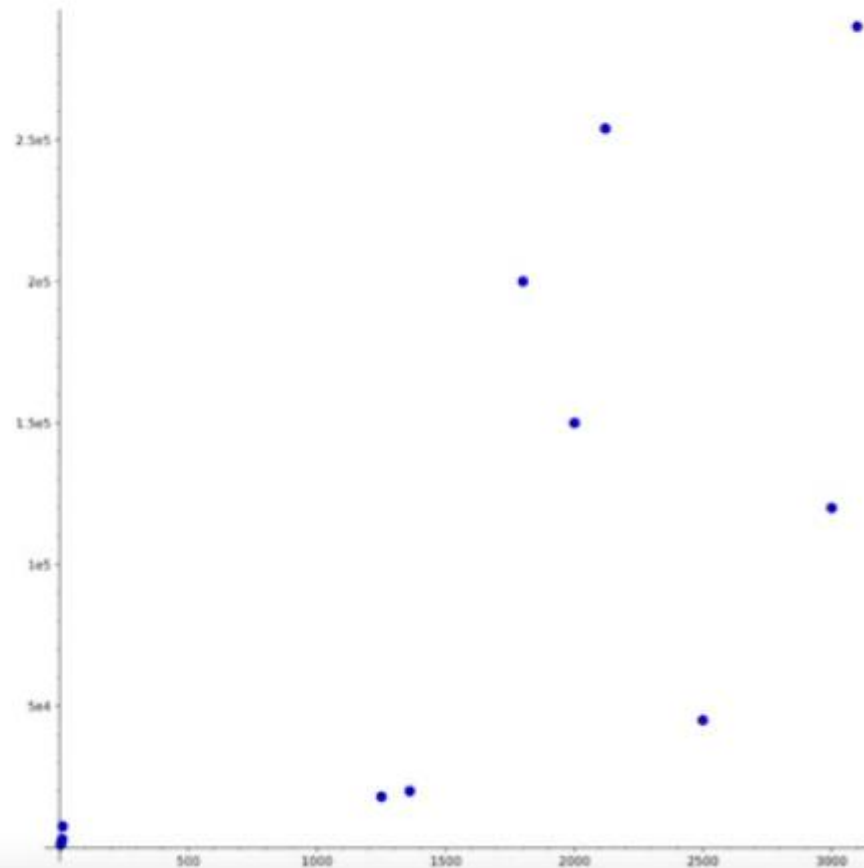
house  = [10, 3000]

apartment  = [1360, 20,000]

office  = [2000, 250,000]

house  = [6, 1650]

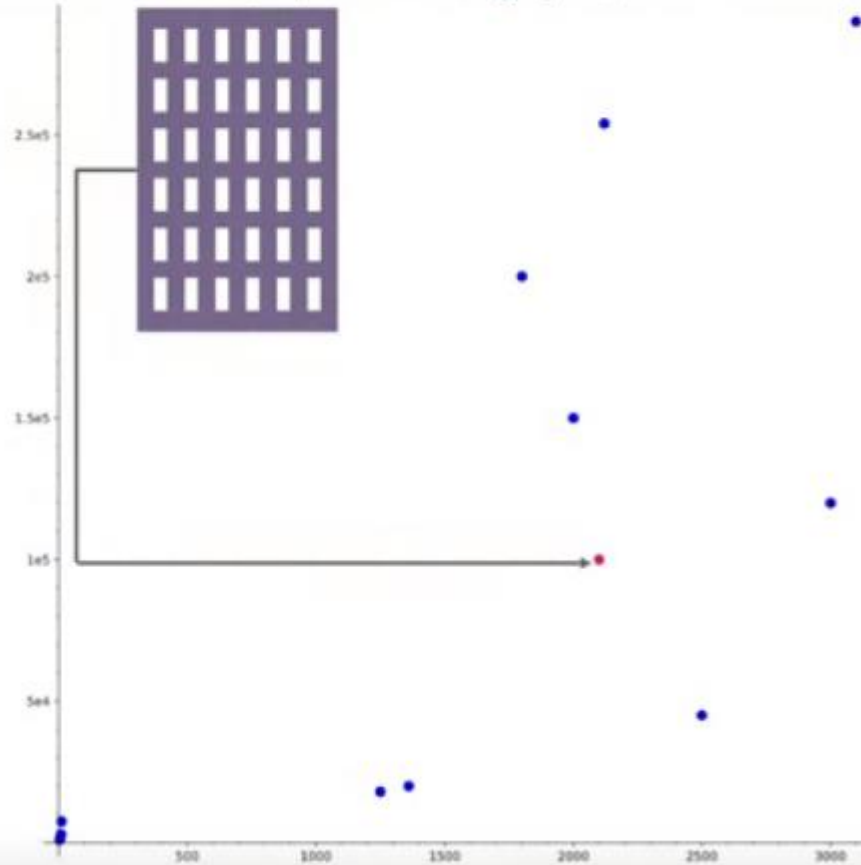
# KNN (K- Nearest Neighbors) Algorithm



10	3000	house
1360	20000	apartment
2000	250000	office
7	2050	house
4	1000	house
3100	290000	office
3000	120000	apartment
12	7500	house
2500	45000	apartment
1800	200000	office
2120	254000	office
⋮	⋮	⋮
1250	18000	apartment

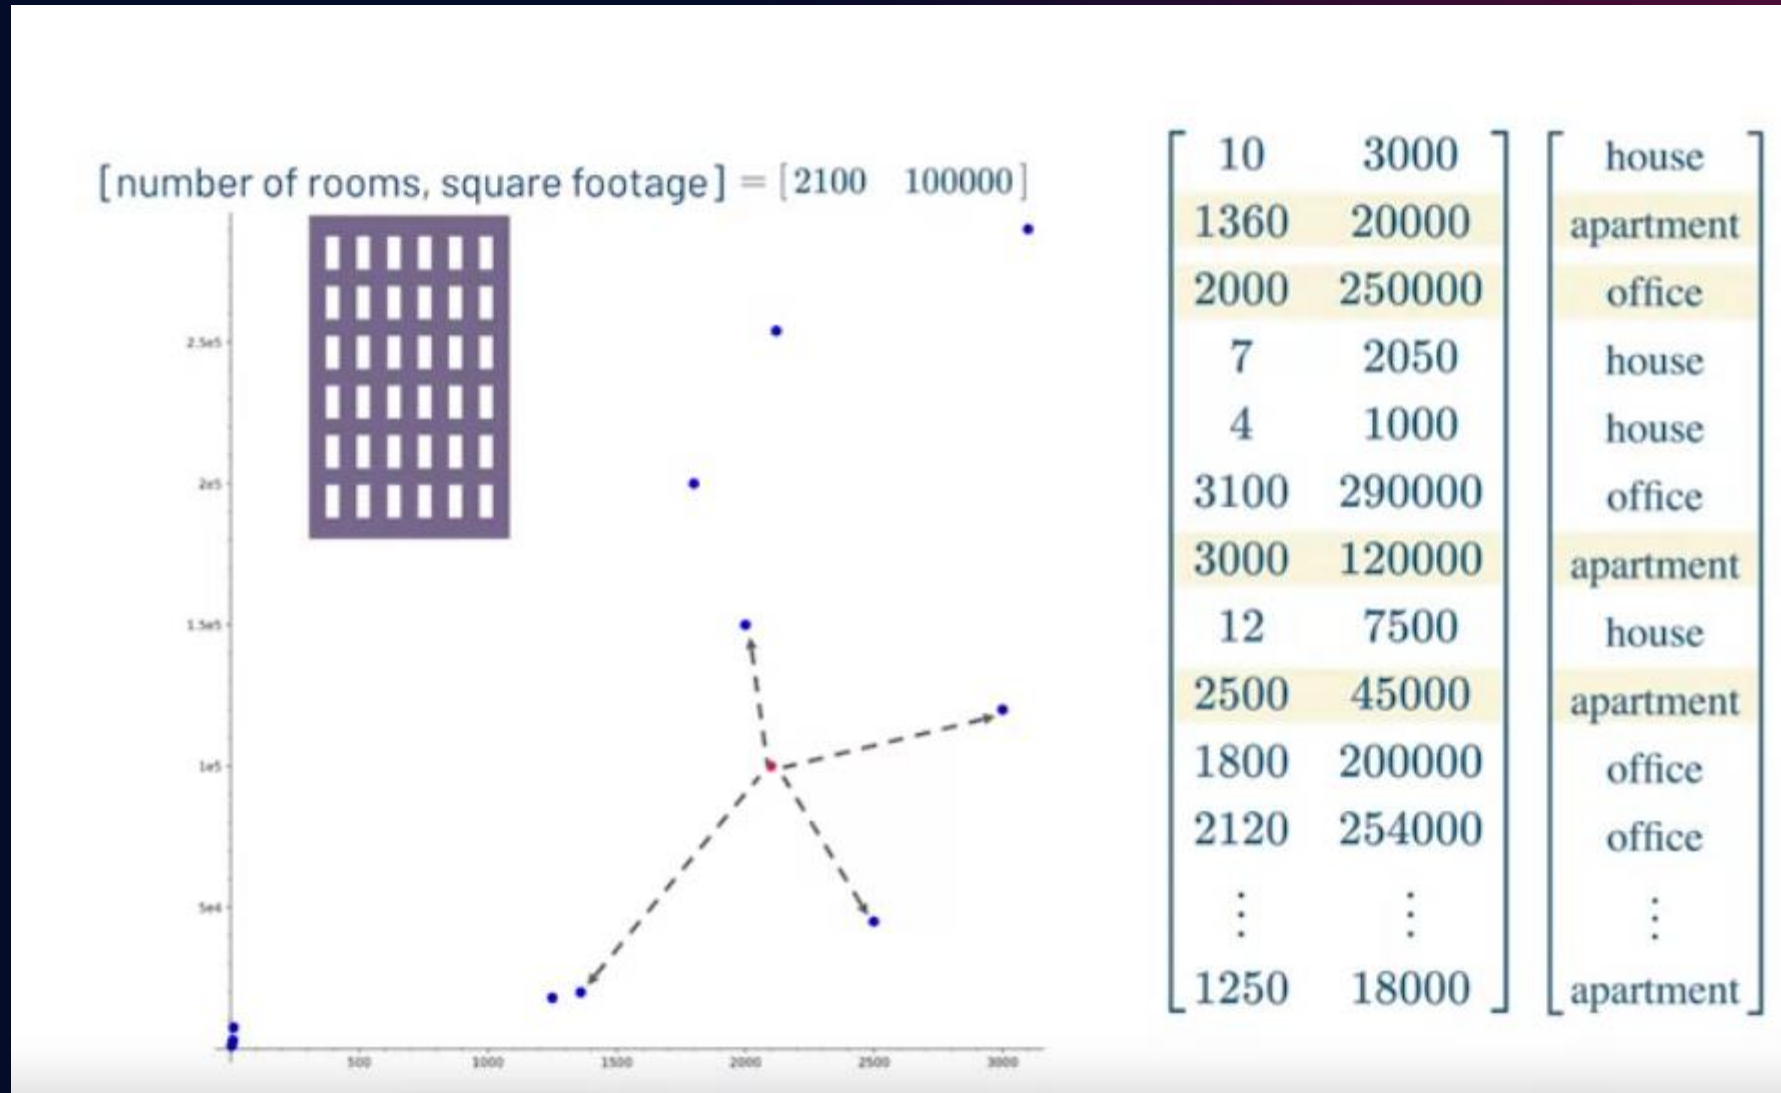
# KNN (K- Nearest Neighbors) Algorithm

[number of rooms, square footage] = [2100 100000]



10	3000	house
1360	20000	apartment
2000	250000	office
7	2050	house
4	1000	house
3100	290000	office
3000	120000	apartment
12	7500	house
2500	45000	apartment
1800	200000	office
2120	254000	office
⋮	⋮	⋮
1250	18000	apartment

# KNN (K- Nearest Neighbors) Algorithm



# KNN (K- Nearest Neighbors) Algorithm

m examples

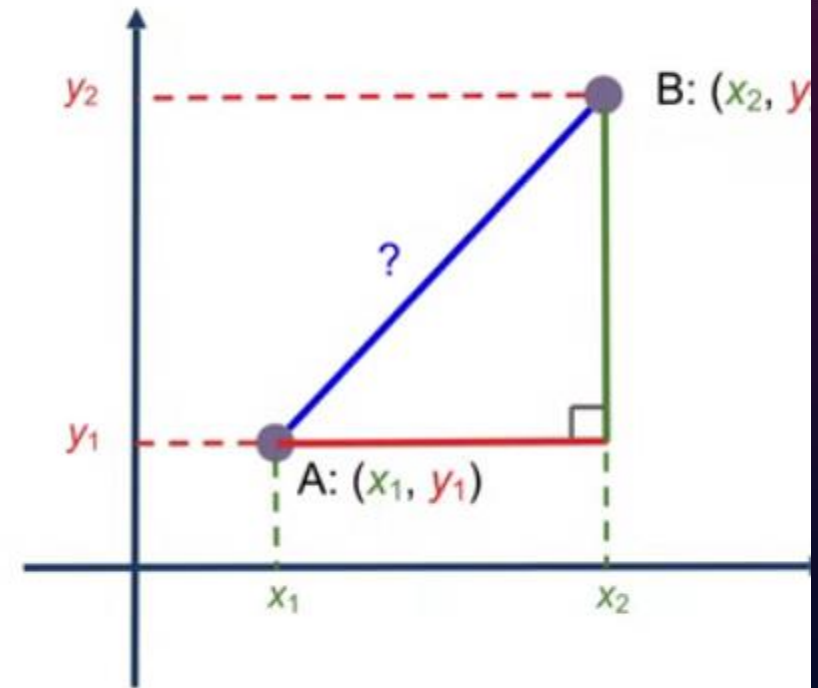
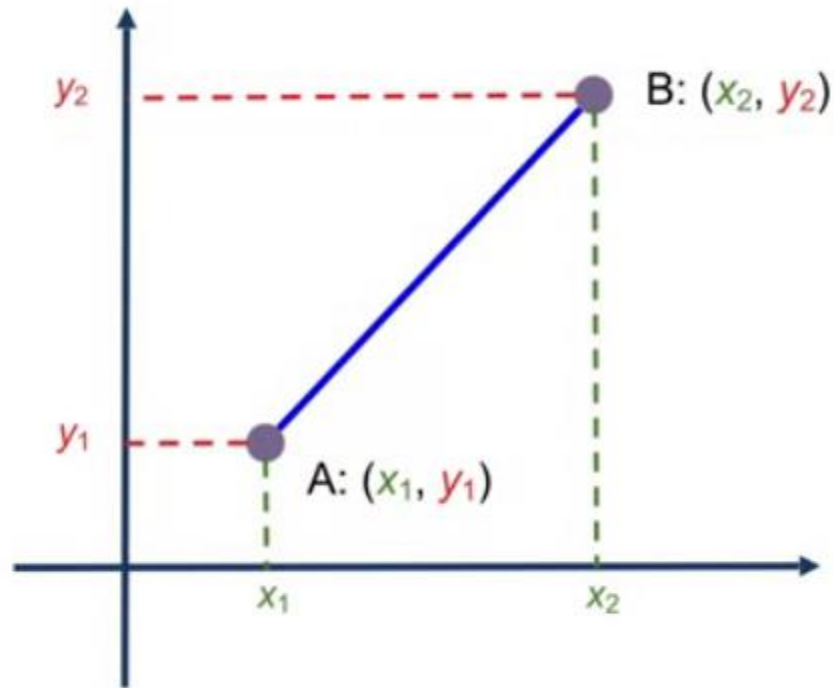
10	3000	150	1	...	0
1360	20000	220	1	...	1
2000	150000	300	0	...	2
7	2050	45	0	...	0
4	1000	13	1	...	0
3100	290000	480	0	...	2
3000	120000	200	1	...	1
12	7500	16	1	...	0
4120	45000	305	0	...	1
1800	200000	122	1	...	2
2120	254000	320	0	...	2
⋮	⋮	⋮	⋮	⋱	⋮
1250	18000	100	0	...	1

n features



# KNN (K- Nearest Neighbors) Algorithm

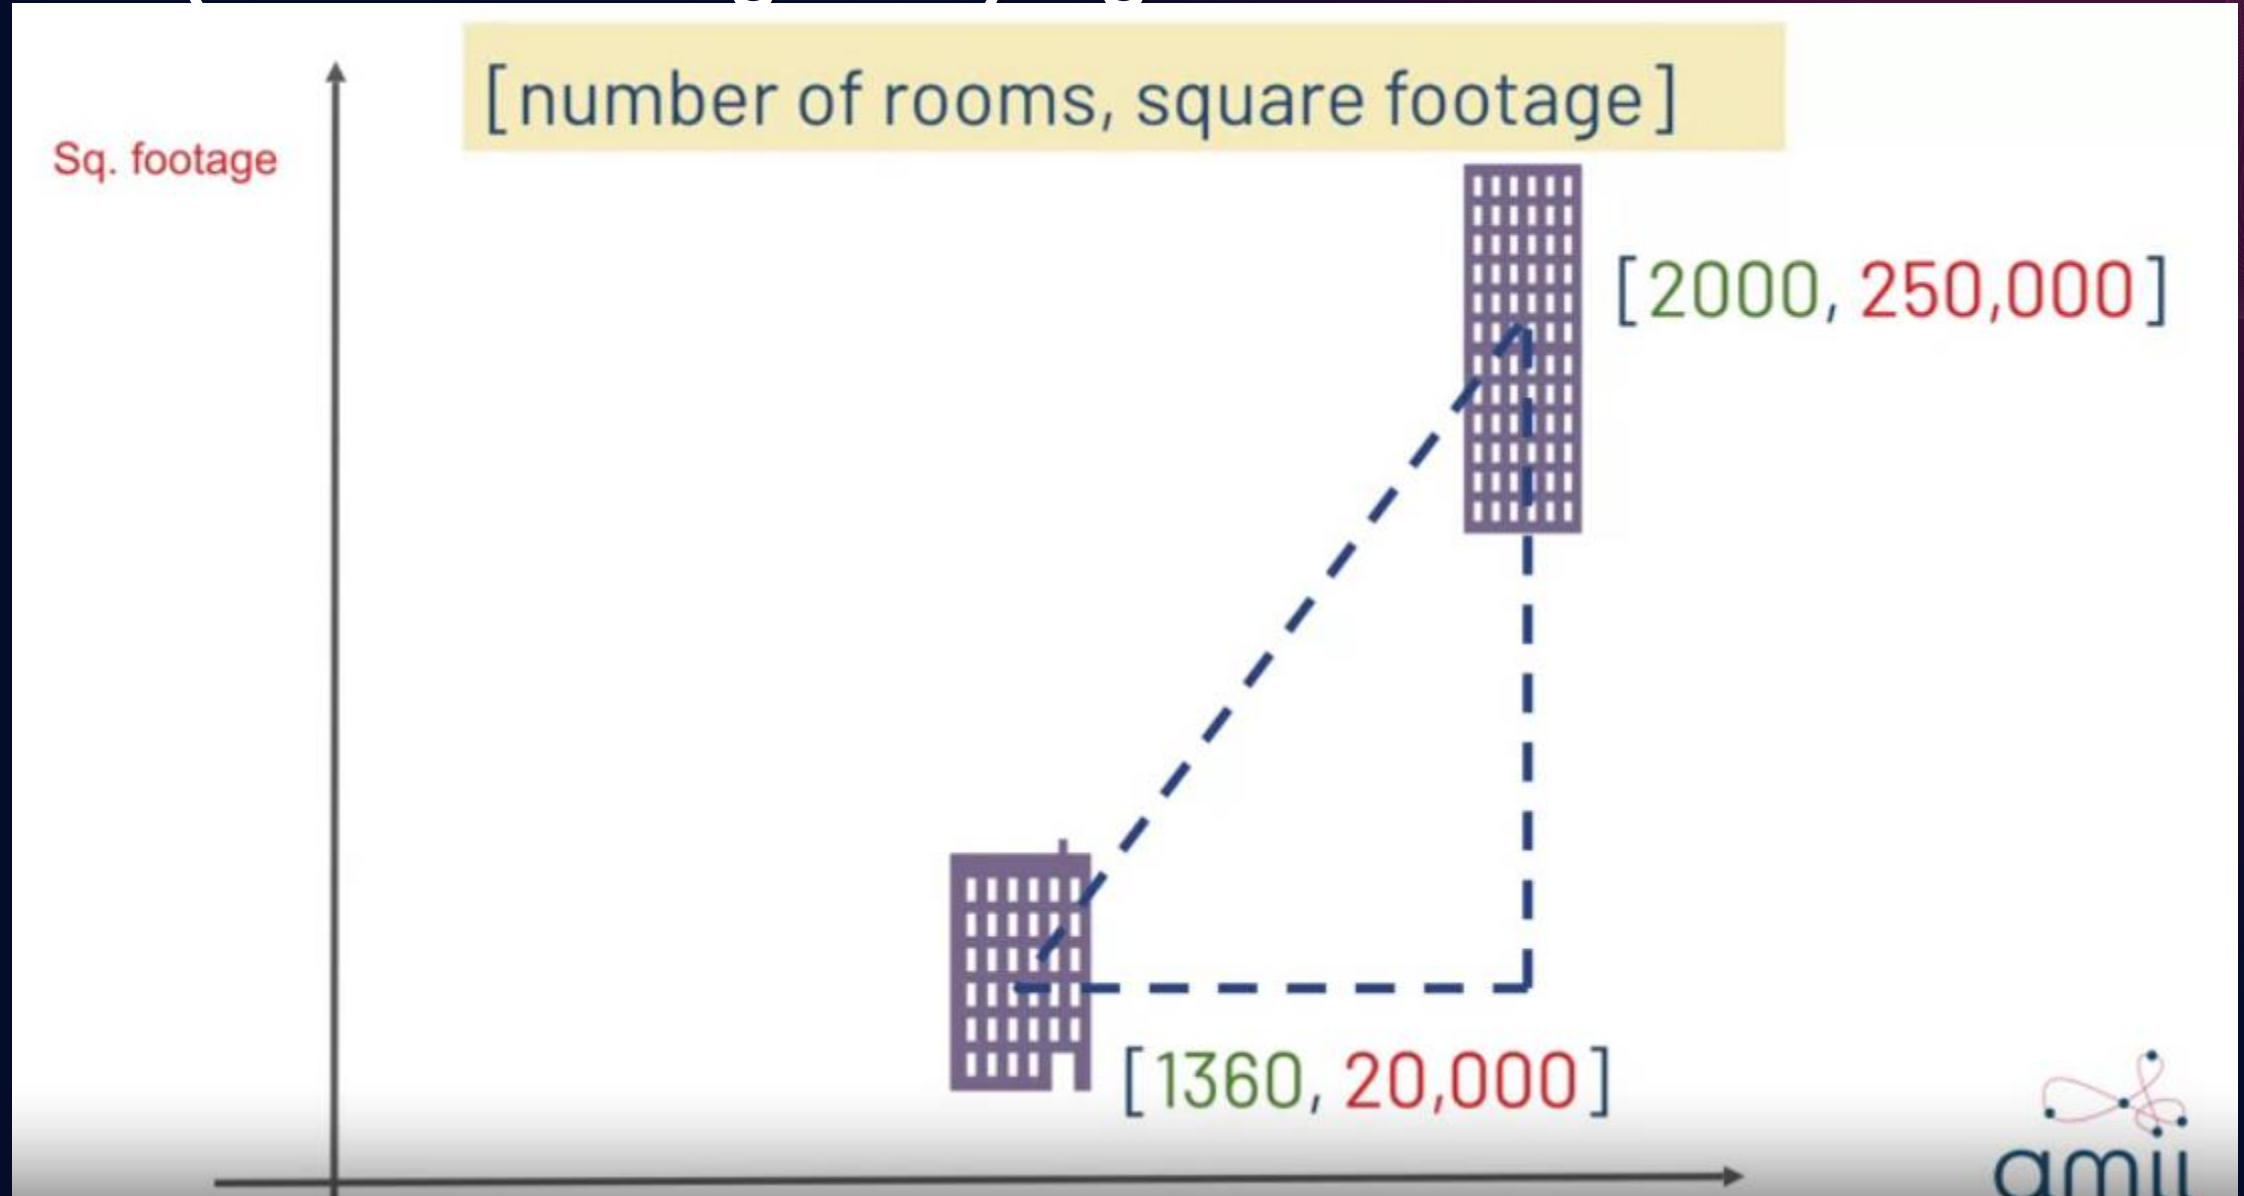
## Pythagorean theorem



$$? = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$



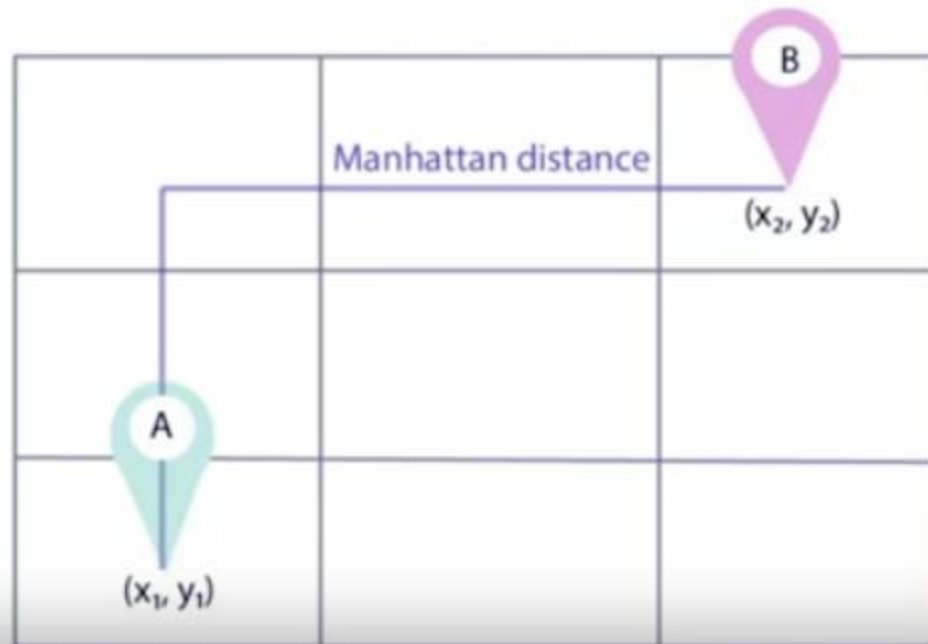
# KNN (K- Nearest Neighbors) Algorithm





# KNN (K- Nearest Neighbors) Algorithm

Manhattan distance

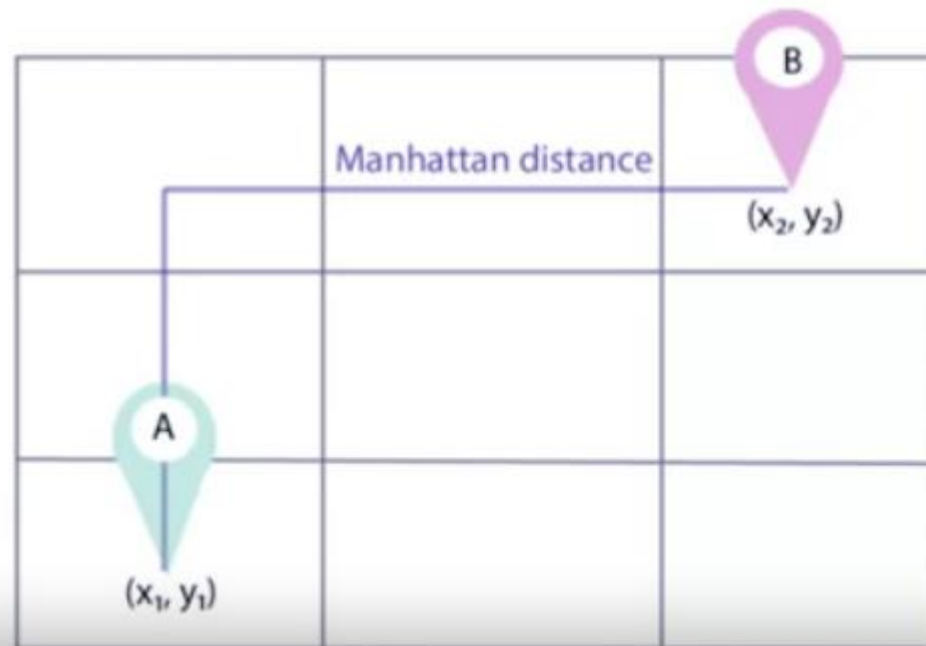




# KNN (K- Nearest Neighbors) Algorithm

Manhattan distance

$$D_M(A, B) = \sum_{j=1}^m |X_{(A,j)} - X_{(B,j)}|$$





# KNN (K- Nearest Neighbors) Algorithm

Hamming distance

$$D_H(A, B) = \sum_{j=1}^m |X_{(A,j)} - X_{(B,j)}|$$

$$X_{(A,j)} = X_{(B,j)} \Rightarrow D_{H_j} = 0$$

$$X_{(A,j)} \neq X_{(B,j)} \Rightarrow D_{H_j} = 1$$