

# Relational Model

---

# Relational Model

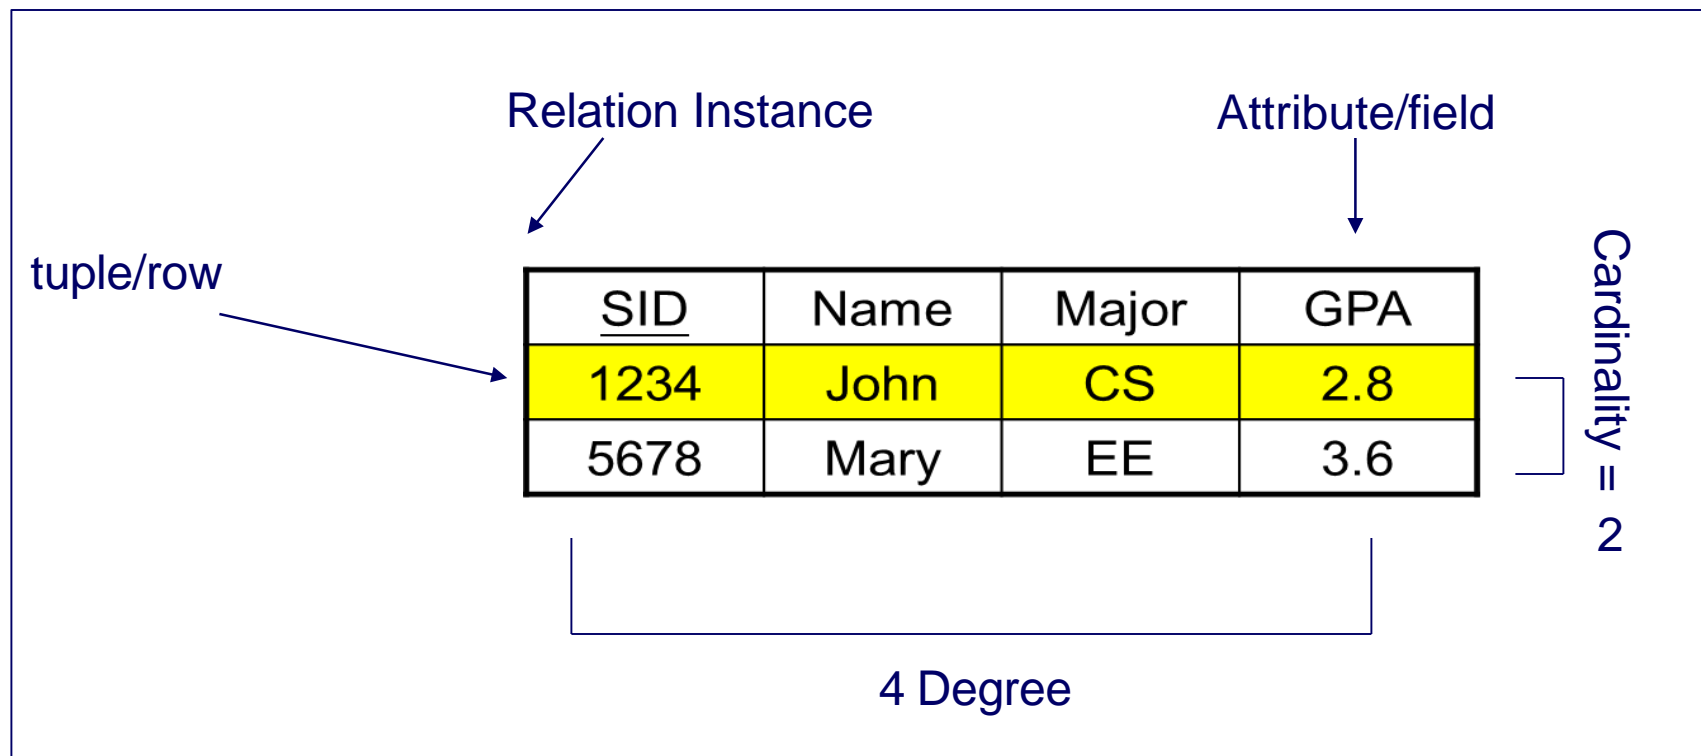
---

- To represent data in the relational mode -> Relation
- A relation consists of
  - Relation Schema
    - Relation's name, Name of each field/column/attribute, domain of each field  
**Student(SID: Integer, Name: String, Major: String, GPA: Number)**
  - Relation Instance
    - Set of tuples which has the same number of fields as the schema

<u>SID</u>	Name	Major	GPA
1234	John	CS	2.8
5678	Mary	EE	3.6

# Relational Model: Requirements

- Rows be unique (key constraint)
- Order of rows does not matter
- Degree of a relation is the number of fields
- Cardinality of a relation is the number of tuples



# Integrity Constraints

---

- Condition specified on a database schema and restrict the data that can be stored in a DB instance.
- Integrity Constraints must be true for any DB instance
- DBMS enforces the integrity constraints.
- Integrity Constraints
  - Domain Constraint
  - Primary Key Constraint
  - Foreign Key Constraint

# Primary Key Constraints

- A set of fields is a key for a relation if :
  1. No two distinct tuples can have same values in all key fields  
  
(SID, Name) -> (1234, John) and (1234, Jack)  
~~(1234, John) and (1234, John)~~
  2. Primary Key cannot be NULL <- Entity Integrity Constraint

# Foreign Keys

---

- Foreign key : Set of fields in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a 'logical pointer'.
- Prevent actions that would destroy links between tables
- Prevent invalid data from being inserted into Foreign Key column.
  - It has to be one of the values contained in the table it points to.
- Foreign Key can be NULL.

# Foreign Keys

---

- An arrow originates from each foreign key and points to the related primary key in the associated relation
- Relations R1 and R2:
  - Attributes in foreign key have the same domain as the primary key of R2
  - Foreign Key is equal to primary key in some rows of the primary table, or else have no value (NULL).

# Referential Integrity

---

- If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.
- A tuple in one relation that refers to another relation must refer to an existing tuple in that relation.
- A referential integrity constraint requires that for each row of a table, the value in the foreign key matches a value in the parent key.



# Enforcing Referential Integrity

<u>actor_id</u>	actor_name
1	Angelina
2	Brad
3	Jennifer

actor_id	movie_id	movie_name
3	1	
4	2	
3	3	

- Referential integrity NOT enforced.

# Enforcing Referential Integrity

- E.g. *sid* is a foreign key referring to **Students**:
  - Enrolled(*sid*: string, *cid*: string, *grade*: string)
  - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.
- What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a Students tuple is deleted?
  - Also delete all Enrolled tuples that refer to it.
  - Disallow deletion of a Students tuple that is referred to.
  - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
  - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting *'unknown'* or *'inapplicable'*.)

# Creating relational tables

---

- Tables are created using CREATE TABLE statements in SQL
- Each attribute is defined, taking the data type and length from the domain definitions
- For example, the attribute Customer\_Name can be defined as a VARCHAR (variable character) type with length 25
- By specifying NOT NULL, each attribute can be constrained from being assigned a null value
- The primary key for each table is specified using the PRIMARY KEY clause at the end of each table definition

# Creating relational tables

---

```
CREATE TABLE CUSTOMER  
  (CUSTOMER_ID  VARCHAR(5)   NOT NULL  
   CUSTOMER_NAME VARCHAR(25) NOT NULL  
   PRIMARY KEY (CUSTOMER_ID);
```

# Creating relational tables

---

```
CREATE TABLE ORDER
  (ORDER_ID      CHAR(5)          NOT NULL
  ORDER_DATE    DATE              NOT NULL
  C_ID VARCHAR(5)                NOT NULL
  PRIMARY KEY (ORDER_ID)
  FOREIGN KEY (CUSTOMER_ID) REFERENCES
    CUSTOMER(CUSTOMER_ID);
```

# Creating relational tables

---

- In SQL, a FOREIGN KEY REFERENCES statement corresponds to one of these arrows
- The foreign key CUSTOMER\_ID references the primary key of CUSTOMER, which is also CUSTOMER\_ID
- Foreign Key constraint has DELETE and UPDATE operations
  - Actions:
    - Cascade, No Action, Set NULL, Set Default

# Creating relational tables

---

- **ON UPDATE/ON DELETE Actions:**
  - **CASCADE:** Delete or update the row from the parent table, and automatically delete or update the matching rows in the child table.
  - **NO ACTION:** Prevent the deletion or update of a parent key if there is a row in the child table that references the key (default)
  - **SET NULL:** Delete or update the row from the parent table, and set the foreign key column or columns in the child table to NULL.
  - **SET DEFAULT:** The "SET DEFAULT" actions are similar to "SET NULL", except that each of the child key columns is set to contain the columns default value instead of NULL.

# Creating relational tables

---

The ORDER\_LINE table illustrates how to specify a primary key when that key is a composite attribute of two foreign keys:

```
CREATE TABLE ORDER_LINE
  (ORDER_ID      CHAR(5)          NOT NULL
   PRODUCT_ID    CHAR(5)          NOT NULL
   QUANTITY      INT              NOT NULL
  PRIMARY KEY (ORDER_ID, PRODUCT_ID)
  FOREIGN KEY (ORDER_ID) REFERENCES ORDER (ORDER_ID)
  FOREIGN KEY (PRODUCT_ID) REFERENCES PRODUCT (PRODUCT_ID);
```



# Constraints in Create Table

---

- Adding constraints to a table enables the database system to enforce data integrity.
- Different types of constraints:

\* Not Null

\* Default Values

\* Unique

\* Primary Key

\* Foreign Key

\* Check Condition

# An Example

---

```
CREATE TABLE Student (  
    ID          NUMBER,  
    Fname       VARCHAR2(20),  
    Lname       VARCHAR2(20),  
);
```

# Not Null Constraint

---

```
CREATE TABLE Student (  
    ID          NUMBER,  
    Fname      VARCHAR2(20) NOT NULL,  
    Lname      VARCHAR2(20) NOT NULL,  
);
```

# Primary Key Constraint

---

```
CREATE TABLE Student (  
    ID          NUMBER PRIMARY KEY,  
    Fname       VARCHAR2(20) NOT NULL,  
    Lname       VARCHAR2(20) NOT NULL,  
);
```

Primary Key implies: \* NOT NULL \* UNIQUE.  
There can only be one primary key.

# Primary Key Constraint (Syntax 2)

---

```
CREATE TABLE Students (  
    ID          NUMBER,  
    Fname       VARCHAR2(20) NOT NULL,  
    Lname       VARCHAR2(20) NOT NULL,  
    PRIMARY KEY(ID)  
);
```

Needed when the primary key is made up of one or more fields

# Logical DB Design: ER to Relational

---


So... how do we convert an ER diagram into a table??

## Basic Ideas:

- Build a table for each entity set
- Build a table for each relationship set (if necessary)
- Make a column in the table for each attribute in the entity set
- Primary Key

# Logical DB Design: ER to Relational


- Entity sets to tables:



```
CREATE TABLE Employees
(  ssn CHAR(11),
   name CHAR(20),
   lot  INTEGER,
   PRIMARY KEY (ssn))
```

<u>ssn</u>	name	lot
123-22-3666	Smith	48
231-31-5368	Lee	22
131-24-3650	Brown	35

# Strong Entity Sets



<u>SSN</u>	Name	Street	City
123-22-3666	Smith	Main St	Santa Clara
231-31-5368	Lee	1st St	San Jose
131-24-3650	Brown	2nd St	San Jose

<u>PID</u>	Name
9999	Smith
4444	Lee
3333	Brown



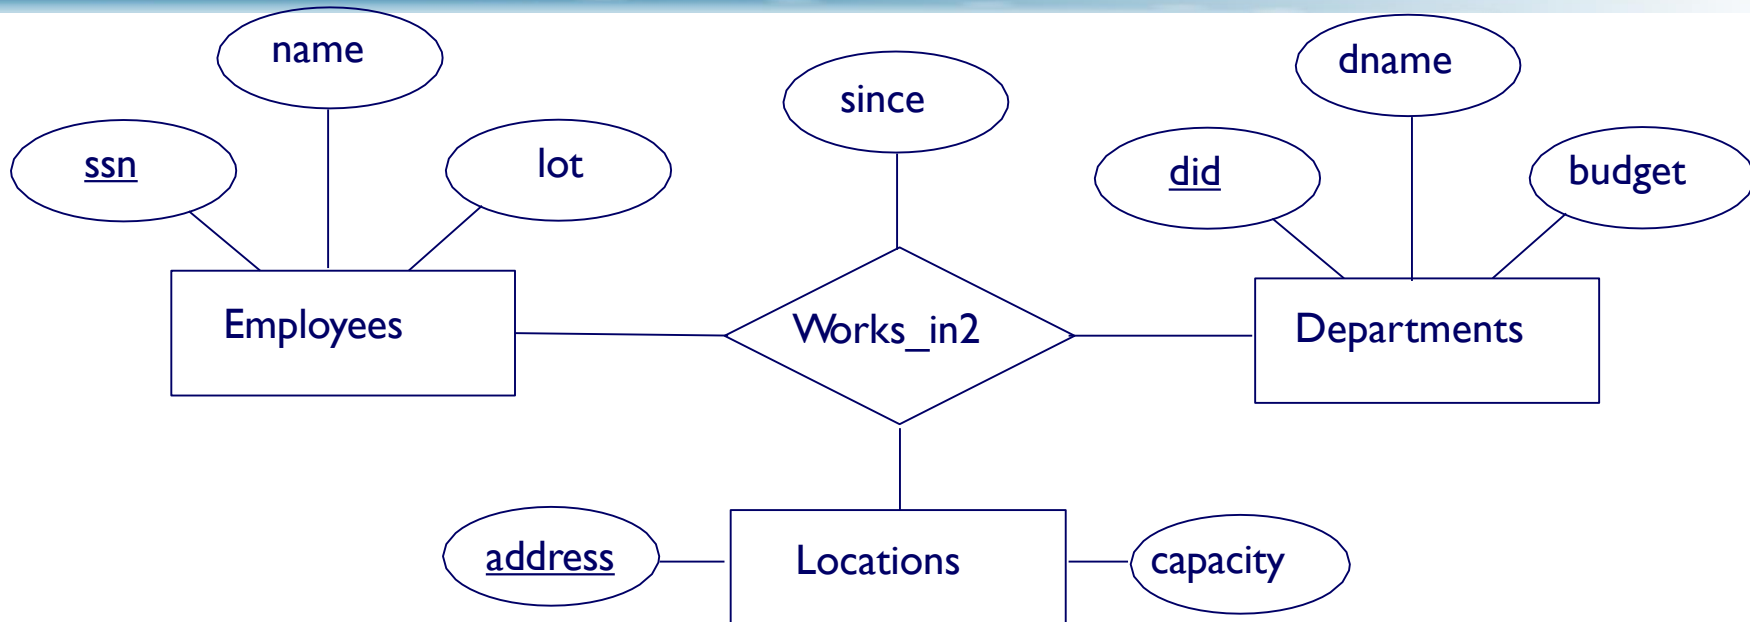
# Relationship Sets to Tables

---

- To represent a relationship, must identify:
  - Each participating entity
    - Keys for each participating entity set (as foreign keys).
    - This set of attributes forms a **superkey** for the relation.
  - All descriptive attributes of the relationship.

# Relationship Sets to Tables

## No Constraints




<u>ssn</u>	<u>did</u>	<u>address</u>	since

- ssn is a foreign key referencing employees
- did is a foreign key referencing departments
- address is a foreign key referencing locations

```
CREATETABLEWorks_In2(  
  ssn CHAR(11),  
  did INTEGER,  
  address CHAR(20),  
  since DATE,  
  PRIMARY KEY (ssn, did, address),  
  FOREIGN KEY (ssn)  
    REFERENCES Employees,  
  FOREIGN KEY (address)  
    REFERENCES Locations,  
  FOREIGN KEY (did)  
    REFERENCES Departments);
```

# Relationship Sets to Tables

## With Key Constraints




- Method I
  - Create a separate relation for Manages
  - did is the key
  - Separate relations for Employees and Departments.

ssn	<u>did</u>	since

```
CREATE TABLE Manages(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    FOREIGN KEY (did) REFERENCES Departments);
```

# Relationship Sets to Tables

## With Key Constraints





<u>did</u>	dname	budget	ssn	since

- Method 2
  - Each department has a unique manager, we could instead combine Manages and Departments.


```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees);
```

# Review: Key Constraints


- Each dept has at most one manager, according to the key constraint on Manages.




**1-to-1**



**1-to Many**



**Many-to-1**




**Many-to-Many**

*Translation to relational model?*

# Review: Participation Constraints

- Does every department have a manager?
  - If so, this is a participation constraint: the participation of Departments in Manages is said to be total (vs. partial).
    - Every did value in Departments table must appear in a row of the Manages table (with a non-null ssn value!)



# Participation Constraints in SQL


---

- We can capture participation constraints involving one entity set in a binary relationship

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE NO ACTION)
```

# Review: Weak Entities

- A weak entity can be identified uniquely only by considering the primary key of another (owner) entity.
  - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
  - Weak entity set must have total participation in this identifying relationship set.





# Translating Weak Entity Sets


---

- Weak Entity Set cannot exist alone
  - When the owner entity is deleted, all owned weak entities must also be deleted.
- To build a table/schema for weak entity set
  - Construct a table with one column for each attribute in the weak entity set
  - Remember to include partial key
  - Include primary key of the Strong Entity Set (the entity set that the weak entity set is depending on) as foreign key
  - Primary Key of the weak entity set = partial key + foreign key

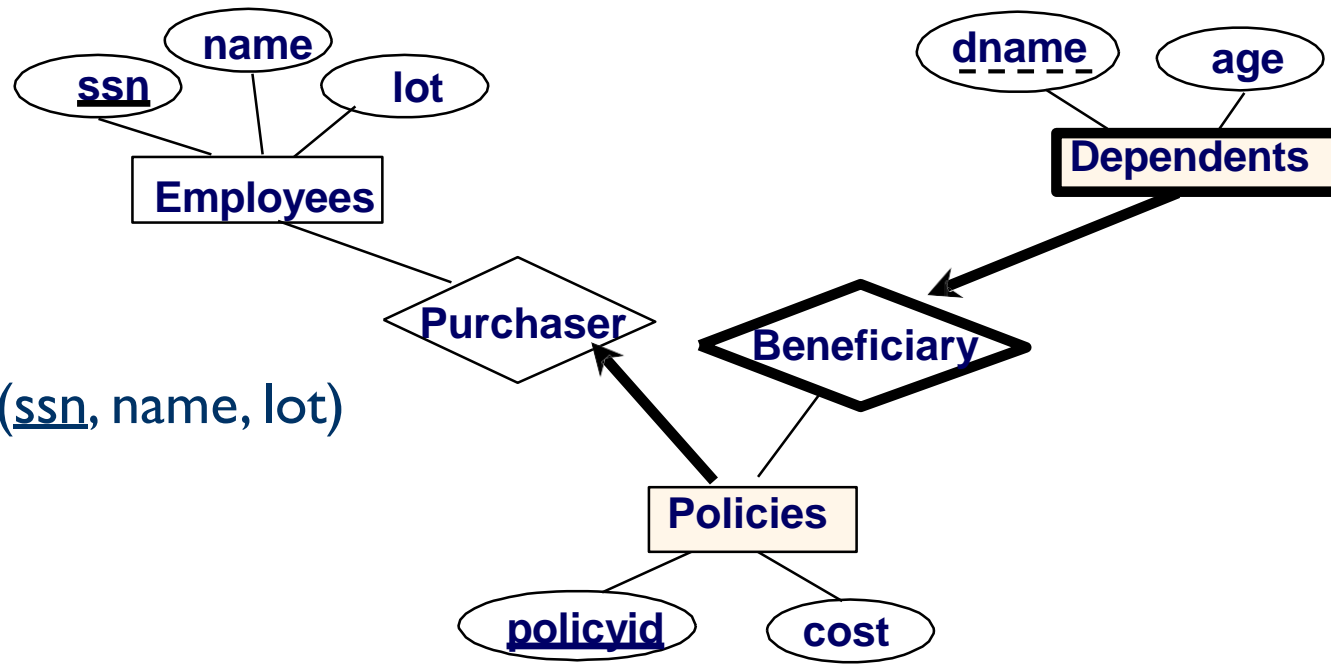
# Translating Weak Entity Sets

- Weak entity set and identifying relationship set are translated into a single table.
  - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
  dname CHAR(20),  
  birthday DATE,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (dname, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees  
  ON DELETE CASCADE);
```



# Review: Binary vs. Ternary



- Employee (ssn, name, lot)
- Dependents(dname, policyid age)
  - policyid is a foreign key referencing Policies
- Policies (policyid, cost, ssn)
  - ssn is a foreign key referencing Employees
  - ssn can not be NULL

# Review: Binary vs. Ternary

---



- The key constraints allow us to combine Purchaser with Policies and Beneficiary with Dependents.
- Participation constraints lead to **NOT NULL** constraints.

```
CREATE TABLE Policies (  
  policyid INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (policyid).  
  FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
  dname CHAR(20),  
  age INTEGER,  
  policyid INTEGER,  
  PRIMARY KEY (dname, policyid).  
  FOREIGN KEY (policyid) REFERENCES Policies,  
    ON DELETE CASCADE)
```

# Representing Composite Attribute

- One column for each component attribute
- NO column for the composite attribute itself



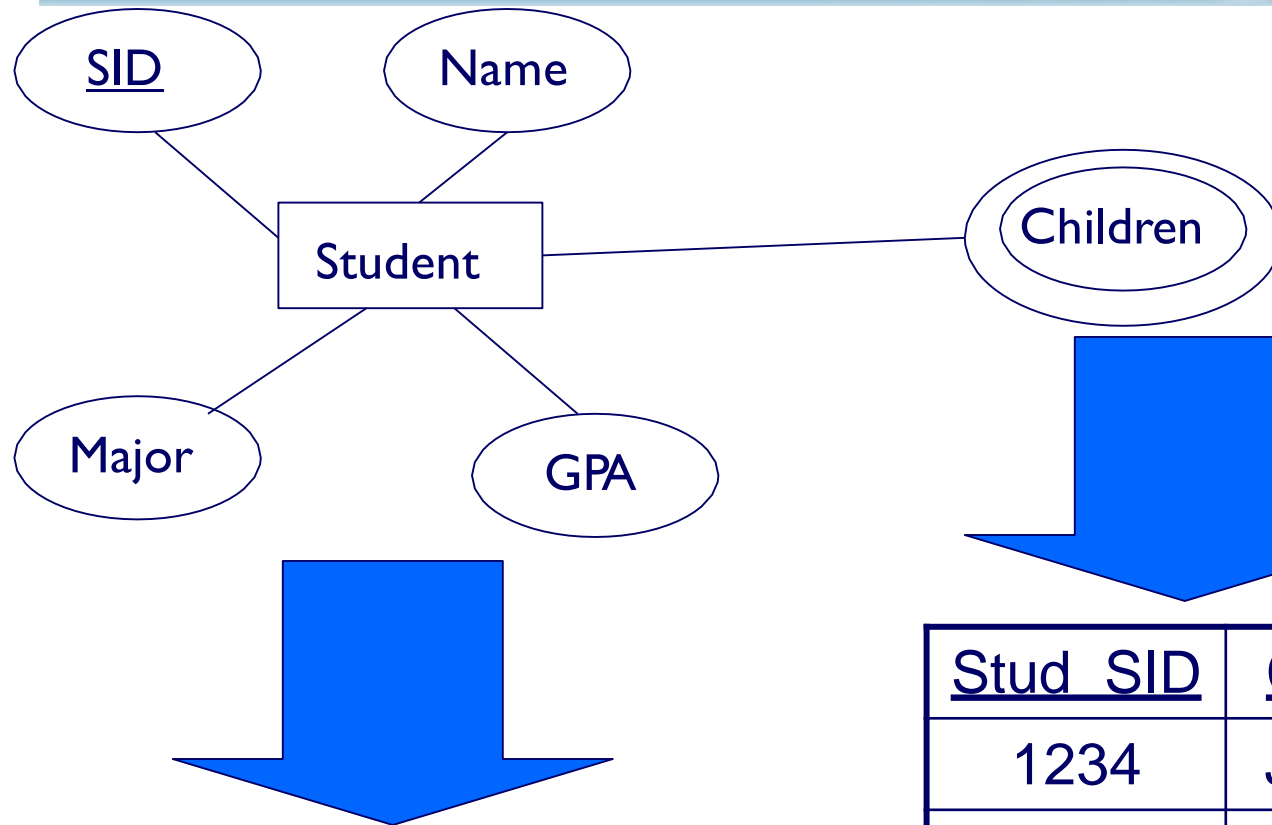
<u>SSN</u>	Name	Street	City
9999	Dr. Smith	50 1 <sup>st</sup> St.	Cupertino
8888	Dr. Lee	1 B St.	San Jose

# Representing Multivalued Attribute

---

- For each multivalued attribute in an entity set/relationship set
  - Build a new relation schema with two columns
  - One column for the primary keys of the entity set/relationship set that has the multivalued attribute
  - Another column for the multivalued attributes. Each cell of this column holds only one value. So each value is represented as a unique tuple
  - Primary key for this schema is the union of all attributes

# Example – Multivalued attribute



<u>SID</u>	Name	Major	GPA
1234	John	CS	2.8
5678	Homer	EE	3.6

The primary key for this table is Student\_SID + Children, the union of all attributes

<u>Stud_SID</u>	<u>Children</u>
1234	Johnson
1234	Mary
5678	Bart
5678	Lisa
5678	Maggie